

```
In [103]: 1 import os  
          2 cwd = os.getcwd()  
          3 cwd
```

```
Out[103]: 'C:\\Users\\USER\\Desktop\\bulldozer-price-prediction'
```

# Predicting the Sale Price of Bulldozers using Machine Learning

In this notebook, we're going to go through an example machine learning project with the goal of predicting the sale price of bulldozers.

## 1. Problem definition

How well can we predict the future sale price of a bulldozer, given its characteristics and previous examples of how much similar bulldozer have been sold for?

## 2. Data

The data is downloaded from the kaggle Bluebook for Bulldozers competition: <https://www.kaggle.com/competitions/bluebook-for-bulldozers/data> (<https://www.kaggle.com/competitions/bluebook-for-bulldozers/data>).

There are 3 main datasets:

- Train.csv is the training set, which contains data through the end of 2011.
- Valid.csv is the validation set, which contains data from January 1, 2012 - April 30, 2012. You make predictions on this set throughout the majority of the competition. Your score on this set is used to create the public leaderboard.
- Test.csv is the test set, which won't be released until the last week of the competition. It contains data from May 1, 2012 - November 2012. Your score on the test set determines your final rank for the competition.

## 3. Evaluation

The evaluation metric for this competition is the RMSLE (root mean squared log error) between the actual and predicted auction prices.

For more on the evaluation of this project check: <https://www.kaggle.com/competitions/bluebook-for-bulldozers/overview/evaluation> (<https://www.kaggle.com/competitions/bluebook-for-bulldozers/overview/evaluation>).

**Note:** The goal for most regression evaluation metrics is to minimize the error. For example, our goal for this project will be to build a machine learning model which minimizes RMSLE.

## 4. Features

Kaggle provides a data dictionary detailing all of the features of the dataset. You can view this data dictionary on excel, on kaggle or Google sheets: <https://www.kaggle.com/competitions/bluebook-for-bulldozers/data> (<https://www.kaggle.com/competitions/bluebook-for-bulldozers/data>).

```
In [104]: 1 import pandas as pd
          2 import numpy as np
          3 import matplotlib.pyplot as plt
          4 import sklearn
```

```
In [105]: 1 # import training and validation sets
          2 df = pd.read_csv('TrainAndValid.csv', low_memory=False)
```

In [106]:

1	<code>df.info()</code>
---	------------------------

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 412698 entries, 0 to 412697
Data columns (total 53 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   SalesID                                   412698 non-null  int64
1   SalePrice                                412698 non-null  float64
2   MachineID                                412698 non-null  int64
3   ModelID                                  412698 non-null  int64
4   datasouce                                412698 non-null  int64
5   auctioneerID                             392562 non-null  float64
6   YearMade                                 412698 non-null  int64
7   MachineHoursCurrentMeter                 147504 non-null  float64
8   UsageBand                                73670 non-null   object
9   saledate                                 412698 non-null  object
10  fiModelDesc                               412698 non-null  object
11  fiBaseModel                               412698 non-null  object
12  fiSecondaryDesc                           271971 non-null  object
13  fiModelSeries                             58667 non-null   object
14  fiModelDescriptor                         74816 non-null   object
15  ProductSize                              196093 non-null  object
16  fiProductClassDesc                       412698 non-null  object
17  state                                    412698 non-null  object
18  ProductGroup                             412698 non-null  object
19  ProductGroupDesc                         412698 non-null  object
20  Drive_System                             107087 non-null  object
21  Enclosure                                 412364 non-null  object
22  Forks                                    197715 non-null  object
23  Pad_Type                                 81096 non-null   object
24  Ride_Control                             152728 non-null  object
25  Stick                                    81096 non-null   object
26  Transmission                             188007 non-null  object
27  Turbocharged                             81096 non-null   object
28  Blade_Extension                           25983 non-null   object
29  Blade_Width                              25983 non-null   object
30  Enclosure_Type                           25983 non-null   object
31  Engine_Horsepower                        25983 non-null   object
32  Hydraulics                               330133 non-null  object
33  Pushblock                               25983 non-null   object
34  Ripper                                   106945 non-null  object
35  Scarifier                                25994 non-null   object
36  Tip_Control                              25983 non-null   object
37  Tire_Size                                97638 non-null   object

```

```
38 Coupler 220679 non-null object
39 Coupler_System 44974 non-null object
40 Grouser_Tracks 44875 non-null object
41 Hydraulics_Flow 44875 non-null object
42 Track_Type 102193 non-null object
43 Undercarriage_Pad_Width 102916 non-null object
44 Stick_Length 102261 non-null object
45 Thumb 102332 non-null object
46 Pattern_Changer 102261 non-null object
47 Grouser_Type 102193 non-null object
48 Backhoe_Mounting 80712 non-null object
49 Blade_Type 81875 non-null object
50 Travel_Controls 81877 non-null object
51 Differential_Type 71564 non-null object
52 Steering_Controls 71522 non-null object
```

```
dtypes: float64(3), int64(5), object(45)
```

```
memory usage: 166.9+ MB
```

```
In [107]: 1 len(df)
```

```
Out[107]: 412698
```

```
In [108]: 1 # to check missing values  
          2 df.isna().sum()
```

```
Out[108]: SalesID          0
          SalePrice        0
          MachineID        0
          ModelID          0
          datasource        0
          auctioneerID     20136
          YearMade         0
          MachineHoursCurrentMeter 265194
          UsageBand        339028
          saledate         0
          fiModelDesc      0
          fiBaseModel      0
          fiSecondaryDesc   140727
          fiModelSeries     354031
          fiModelDescriptor 337882
          ProductSize       216605
          fiProductClassDesc 0
          state            0
          ProductGroup      0
          ProductGroupDesc  0
          Drive_System      305611
          Enclosure         334
          Forks             214983
          Pad_Type          331602
          Ride_Control      259970
          Stick             331602
          Transmission      224691
          Turbocharged      331602
          Blade_Extension   386715
          Blade_Width       386715
          Enclosure_Type    386715
          Engine_Horsepower 386715
          Hydraulics        82565
          Pushblock         386715
          Ripper            305753
          Scarifier         386704
          Tip_Control       386715
          Tire_Size         315060
          Coupler           192019
          Coupler_System    367724
          Grouser_Tracks    367823
          Hydraulics_Flow   367823
          Track_Type        310505
```


```
Undercarriage_Pad_Width    309782
Stick_Length               310437
Thumb                     310366
Pattern_Changer            310437
Grouser_Type               310505
Backhoe_Mounting           331986
Blade_Type                 330823
Travel_Controls            330821
Differential_Type          341134
Steering_Controls          341176
dtype: int64
```

```
In [109]: 1 # to find the column names
          2 df.columns
```

```
Out[109]: Index(['SalesID', 'SalePrice', 'MachineID', 'ModelID', 'datasource',
                 'auctioneerID', 'YearMade', 'MachineHoursCurrentMeter', 'UsageBand',
                 'saledate', 'fiModelDesc', 'fiBaseModel', 'fiSecondaryDesc',
                 'fiModelSeries', 'fiModelDescriptor', 'ProductSize',
                 'fiProductClassDesc', 'state', 'ProductGroup', 'ProductGroupDesc',
                 'Drive_System', 'Enclosure', 'Forks', 'Pad_Type', 'Ride_Control',
                 'Stick', 'Transmission', 'Turbocharged', 'Blade_Extension',
                 'Blade_Width', 'Enclosure_Type', 'Engine_Horsepower', 'Hydraulics',
                 'Pushblock', 'Ripper', 'Scarifier', 'Tip_Control', 'Tire_Size',
                 'Coupler', 'Coupler_System', 'Grouser_Tracks', 'Hydraulics_Flow',
                 'Track_Type', 'Undercarriage_Pad_Width', 'Stick_Length', 'Thumb',
                 'Pattern_Changer', 'Grouser_Type', 'Backhoe_Mounting', 'Blade_Type',
                 'Travel_Controls', 'Differential_Type', 'Steering_Controls'],
                 dtype='object')
```



```
In [110]: 1 fig, ax = plt.subplots()
          2 ax.scatter(df['saledate'][:1000], df['SalePrice'][:1000]);
```




```
In [111]: 1 df.saledate[:1000]
```

```
Out[111]: 0      11/16/2006 0:00
          1      3/26/2004 0:00
          2      2/26/2004 0:00
          3      5/19/2011 0:00
          4      7/23/2009 0:00
          ...
          995    7/16/2009 0:00
          996    6/14/2007 0:00
          997    9/22/2005 0:00
          998    7/28/2005 0:00
          999    6/16/2011 0:00
          Name: saledate, Length: 1000, dtype: object
```

```
In [112]: 1 df.saledate.dtypes
```

```
Out[112]: dtype('O')
```

```
In [113]: 1 df.SalePrice.plot.hist();
```



## Parsing dates

When we work with time series data, we want to enrich the time & date component as much as possible.

we can do that by telling pandas which of our columns has date in it using the `parse_date` parameter.

```
In [114]: 1 # import data again but this time parse dates
          2
          3 df = pd.read_csv('TrainAndValid.csv',
          4                     low_memory=False,
          5                     parse_dates=['saledate'])
```


```
In [115]: 1 df.saledate.dtypes
```

```
Out[115]: dtype('<M8[ns]')
```

```
In [116]: 1 df.saledate[: 1000]
```

```
Out[116]: 0      2006-11-16
          1      2004-03-26
          2      2004-02-26
          3      2011-05-19
          4      2009-07-23
          ..
          995    2009-07-16
          996    2007-06-14
          997    2005-09-22
          998    2005-07-28
          999    2011-06-16
          Name: saledate, Length: 1000, dtype: datetime64[ns]
```

```
In [117]: 1 fig, ax = plt.subplots()
          2 ax.scatter(df['saledate'][: 1000], df['SalePrice'][: 1000]);
```



In [118]:

```
1 df.head()
```

Out[118]:

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHours	CurrentMeter	UsageBand	saledate	...	l
0	1139246	66000.0	999089	3157	121	3.0	2004		68.0	Low	2006-11-16	...	
1	1139248	57000.0	117657	77	121	3.0	1996		4640.0	Low	2004-03-26	...	
2	1139249	10000.0	434808	7009	121	3.0	2001		2838.0	High	2004-02-26	...	
3	1139251	38500.0	1026470	332	121	3.0	2001		3486.0	High	2011-05-19	...	
4	1139253	11000.0	1057373	17311	121	3.0	2007		722.0	Medium	2009-07-23	...	

5 rows × 53 columns

In [119]:

1	<code>df.head().T</code>
---	--------------------------

Out[119]:

	0	1	2	3	4
<b>SalesID</b>	1139246	1139248	1139249	1139251	1139253
<b>SalePrice</b>	66000.0	57000.0	10000.0	38500.0	11000.0
<b>MachineID</b>	999089	117657	434808	1026470	1057373
<b>ModelID</b>	3157	77	7009	332	17311
<b>datasource</b>	121	121	121	121	121
<b>auctioneerID</b>	3.0	3.0	3.0	3.0	3.0
<b>YearMade</b>	2004	1996	2001	2001	2007
<b>MachineHoursCurrentMeter</b>	68.0	4640.0	2838.0	3486.0	722.0
<b>UsageBand</b>	Low	Low	High	High	Medium
<b>saledate</b>	2006-11-16 00:00:00	2004-03-26 00:00:00	2004-02-26 00:00:00	2011-05-19 00:00:00	2009-07-23 00:00:00
<b>fiModelDesc</b>	521D	950FII	226	PC120-6E	S175
<b>fiBaseModel</b>	521	950	226	PC120	S175
<b>fiSecondaryDesc</b>	D	F	NaN	NaN	NaN
<b>fiModelSeries</b>	NaN	II	NaN	-6E	NaN
<b>fiModelDescriptor</b>	NaN	NaN	NaN	NaN	NaN
<b>ProductSize</b>	NaN	Medium	NaN	Small	NaN
<b>fiProductClassDesc</b>	Wheel Loader - 110.0 to 120.0 Horsepower	Wheel Loader - 150.0 to 175.0 Horsepower	Skid Steer Loader - 1351.0 to 1601.0 Lb Operat...	Hydraulic Excavator, Track - 12.0 to 14.0 Metr...	Skid Steer Loader - 1601.0 to 1751.0 Lb Operat...
<b>state</b>	Alabama	North Carolina	New York	Texas	New York
<b>ProductGroup</b>	WL	WL	SSL	TEX	SSL
<b>ProductGroupDesc</b>	Wheel Loader	Wheel Loader	Skid Steer Loaders	Track Excavators	Skid Steer Loaders
<b>Drive_System</b>	NaN	NaN	NaN	NaN	NaN
<b>Enclosure</b>	EROPS w AC	EROPS w AC	OROPS	EROPS w AC	EROPS
<b>Forks</b>	None or Unspecified	None or Unspecified	None or Unspecified	NaN	None or Unspecified
<b>Pad_Type</b>	NaN	NaN	NaN	NaN	NaN

	0	1	2	3	4
<b>Ride_Control</b>	None or Unspecified	None or Unspecified	NaN	NaN	NaN
<b>Stick</b>	NaN	NaN	NaN	NaN	NaN
<b>Transmission</b>	NaN	NaN	NaN	NaN	NaN
<b>Turbocharged</b>	NaN	NaN	NaN	NaN	NaN
<b>Blade_Extension</b>	NaN	NaN	NaN	NaN	NaN
<b>Blade_Width</b>	NaN	NaN	NaN	NaN	NaN
<b>Enclosure_Type</b>	NaN	NaN	NaN	NaN	NaN
<b>Engine_Horsepower</b>	NaN	NaN	NaN	NaN	NaN
<b>Hydraulics</b>	2 Valve	2 Valve	Auxiliary	2 Valve	Auxiliary
<b>Pushblock</b>	NaN	NaN	NaN	NaN	NaN
<b>Ripper</b>	NaN	NaN	NaN	NaN	NaN
<b>Scarifier</b>	NaN	NaN	NaN	NaN	NaN
<b>Tip_Control</b>	NaN	NaN	NaN	NaN	NaN
<b>Tire_Size</b>	None or Unspecified	23.5	NaN	NaN	NaN
<b>Coupler</b>	None or Unspecified	None or Unspecified	None or Unspecified	None or Unspecified	None or Unspecified
<b>Coupler_System</b>	NaN	NaN	None or Unspecified	NaN	None or Unspecified
<b>Grouser_Tracks</b>	NaN	NaN	None or Unspecified	NaN	None or Unspecified
<b>Hydraulics_Flow</b>	NaN	NaN	Standard	NaN	Standard
<b>Track_Type</b>	NaN	NaN	NaN	NaN	NaN
<b>Undercarriage_Pad_Width</b>	NaN	NaN	NaN	NaN	NaN
<b>Stick_Length</b>	NaN	NaN	NaN	NaN	NaN
<b>Thumb</b>	NaN	NaN	NaN	NaN	NaN
<b>Pattern_Changer</b>	NaN	NaN	NaN	NaN	NaN
<b>Grouser_Type</b>	NaN	NaN	NaN	NaN	NaN
<b>Backhoe_Mounting</b>	NaN	NaN	NaN	NaN	NaN
<b>Blade_Type</b>	NaN	NaN	NaN	NaN	NaN



	0	1	2	3	4
<b>Travel_Controls</b>	NaN	NaN	NaN	NaN	NaN
<b>Differential_Type</b>	Standard	Standard	NaN	NaN	NaN
<b>Steering_Controls</b>	Conventional	Conventional	NaN	NaN	NaN

In [120]: 1 df.saledate.head(20)

Out[120]:

0	2006-11-16
1	2004-03-26
2	2004-02-26
3	2011-05-19
4	2009-07-23
5	2008-12-18
6	2004-08-26
7	2005-11-17
8	2009-08-27
9	2007-08-09
10	2008-08-21
11	2006-08-24
12	2005-10-20
13	2006-01-26
14	2006-01-03
15	2006-11-16
16	2007-06-14
17	2010-01-28
18	2006-03-09
19	2005-11-17

Name: saledate, dtype: datetime64[ns]

## Sort DataFrame by saledate

when working with time series data, it's a good idea to sort it by date.

```
In [121]: 1 # sort DataFrame in date order
          2 df.sort_values(by=['saledate'], inplace=True, ascending=True)
          3 df.saledate.head(20)
```

```
Out[121]: 205615    1989-01-17
          274835    1989-01-31
          141296    1989-01-31
          212552    1989-01-31
          62755     1989-01-31
          54653     1989-01-31
          81383     1989-01-31
          204924    1989-01-31
          135376    1989-01-31
          113390    1989-01-31
          113394    1989-01-31
          116419    1989-01-31
          32138     1989-01-31
          127610    1989-01-31
          76171     1989-01-31
          127000    1989-01-31
          128130    1989-01-31
          127626    1989-01-31
          55455     1989-01-31
          55454     1989-01-31
          Name: saledate, dtype: datetime64[ns]
```

## make a copy of the original DataFrame

We make a copy of the original dataframe so when we manipulate the copy, we've still got our original data.

```
In [122]: 1 # make a copy
          2 df_tmp = df.copy()
```

```
In [123]: 1 df_tmp.saledate.head(20)
```

```
Out[123]: 205615    1989-01-17
274835    1989-01-31
141296    1989-01-31
212552    1989-01-31
62755     1989-01-31
54653     1989-01-31
81383     1989-01-31
204924    1989-01-31
135376    1989-01-31
113390    1989-01-31
113394    1989-01-31
116419    1989-01-31
32138     1989-01-31
127610    1989-01-31
76171     1989-01-31
127000    1989-01-31
128130    1989-01-31
127626    1989-01-31
55455     1989-01-31
55454     1989-01-31
Name: saledate, dtype: datetime64[ns]
```

## Feature Engineering

Add datetime parameters for saledate column

```
In [124]: 1 df_tmp[: 1].saledate.dt.year
```

```
Out[124]: 205615    1989
Name: saledate, dtype: int64
```

```
In [125]: 1 df_tmp[: 1].saledate.dt.day
```

```
Out[125]: 205615    17
Name: saledate, dtype: int64
```

In [126]: 1 df\_tmp[: 1].saledate.dt.month

Out[126]: 205615 1  
Name: saledate, dtype: int64

In [127]: 1 df\_tmp[:1].saledate

Out[127]: 205615 1989-01-17  
Name: saledate, dtype: datetime64[ns]

In [128]: 1 df\_tmp['saleYear'] = df\_tmp.saledate.dt.year  
2 df\_tmp['saleMonth'] = df\_tmp.saledate.dt.month  
3 df\_tmp['saleDay'] = df\_tmp.saledate.dt.day  
4 df\_tmp['saleDayOfWeek'] = df\_tmp.saledate.dt.dayofweek  
5 df\_tmp['saleDayOfYear'] = df\_tmp.saledate.dt.dayofyear

In [129]: 1 df\_tmp.head().T

<b>Pattern_Changer</b>	NaN	NaN	NaN	NaN	NaN
<b>Grouser_Type</b>	NaN	NaN	NaN	NaN	NaN
<b>Backhoe_Mounting</b>	None or Unspecified	NaN	None or Unspecified	NaN	None or Unspecified
<b>Blade_Type</b>	Straight	NaN	Straight	NaN	PAT
<b>Travel_Controls</b>	None or Unspecified	NaN	None or Unspecified	NaN	Lever
<b>Differential_Type</b>	NaN	Standard	NaN	Standard	NaN
<b>Steering_Controls</b>	NaN	Conventional	NaN	Conventional	NaN
<b>saleYear</b>	1989	1989	1989	1989	1989
<b>saleMonth</b>	1	1	1	1	1
<b>saleDay</b>	17	31	31	31	31
<b>saleDayOfWeek</b>	1	1	1	1	1
<b>saleDayOfYear</b>	17	31	31	31	31

```
In [130]: 1 # Now we've enrich our DataFrame with date time features, we can remove 'saledate'  
          2 df_tmp.drop('saledate', axis = 1, inplace=True)
```

```
In [131]: 1 # check the values of different columns  
          2 df_tmp.state.value_counts()
```

```
Out[131]: Florida      67320
          Texas       53110
          California   29761
          Washington   16222
          Georgia      14633
          Maryland     13322
          Mississippi  13240
          Ohio         12369
          Illinois     11540
          Colorado     11529
          New Jersey   11156
          North Carolina 10636
          Tennessee   10298
          Alabama      10292
          Pennsylvania 10234
          South Carolina 9951
          Arizona      9364
          New York     8639
          Connecticut  8276
          Minnesota    7885
          Missouri     7178
          Nevada       6932
          Louisiana    6627
          Kentucky     5351
          Maine        5096
          Indiana      4124
          Arkansas     3933
          New Mexico   3631
          Utah         3046
          Unspecified  2801
          Wisconsin    2745
          New Hampshire 2738
          Virginia     2353
          Idaho        2025
          Oregon       1911
          Michigan     1831
          Wyoming      1672
          Montana      1336
          Iowa         1336
          Oklahoma     1326
          Nebraska     866
          West Virginia 840
          Kansas       667
```

```

Delaware          510
North Dakota      480
Alaska            430
Massachusetts     347
Vermont           300
South Dakota      244
Hawaii            118
Rhode Island      83
Puerto Rico       42
Washington DC     2
Name: state, dtype: int64

```

## 5. Modelling

we've done enough EDA (we could always do more) but let's start to do some model-driven EDA.


In [132]:

```
1 df_tmp.head()
```

Out[132]:

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter	UsageBand	fiModelID
<b>205615</b>	1646770	9500.0	1126363	8434	132	18.0	1974	NaN	NaN	TI
<b>274835</b>	1821514	14000.0	1194089	10150	132	99.0	1980	NaN	NaN	
<b>141296</b>	1505138	50000.0	1473654	4139	132	99.0	1978	NaN	NaN	[
<b>212552</b>	1671174	16000.0	1327630	8591	132	99.0	1980	NaN	NaN	
<b>62755</b>	1329056	22000.0	1336053	4089	132	99.0	1984	NaN	NaN	[

5 rows × 57 columns





In [133]:

```
1 # Let's build a machine Learning model
2 from sklearn.ensemble import RandomForestRegressor
3
4 model = RandomForestRegressor(n_jobs=-1,
5                               random_state=42)
6 model.fit(df_tmp.drop('SalePrice', axis = 1), df_tmp['SalePrice'])
```

```

-----
ValueError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_6476\4172265634.py in <module>
      4 model = RandomForestRegressor(n_jobs=-1,
      5                               random_state=42)
----> 6 model.fit(df_tmp.drop('SalePrice', axis = 1), df_tmp['SalePrice'])

~\anaconda3\lib\site-packages\sklearn\ensemble\_forest.py in fit(self, X, y, sample_weight)
    325         if issparse(y):
    326             raise ValueError("sparse multilabel-indicator for y is not supported.")
--> 327         X, y = self._validate_data(
    328             X, y, multi_output=True, accept_sparse="csc", dtype=DTYPE
    329         )

~\anaconda3\lib\site-packages\sklearn\base.py in _validate_data(self, X, y, reset, validate_separately,
**check_params)
    579         y = check_array(y, **check_y_params)
    580     else:
--> 581         X, y = check_X_y(X, y, **check_params)
    582         out = X, y
    583

~\anaconda3\lib\site-packages\sklearn\utils\validation.py in check_X_y(X, y, accept_sparse, accept_large
_sparse, dtype, order, copy, force_all_finite, ensure_2d, allow_nd, multi_output, ensure_min_samples, en
sure_min_features, y_numeric, estimator)
    962         raise ValueError("y cannot be None")
    963
--> 964     X = check_array(
    965         X,
    966         accept_sparse=accept_sparse,

~\anaconda3\lib\site-packages\sklearn\utils\validation.py in check_array(array, accept_sparse, accept_la
rge_sparse, dtype, order, copy, force_all_finite, ensure_2d, allow_nd, ensure_min_samples, ensure_min_fe
atures, estimator)
    744         array = array.astype(dtype, casting="unsafe", copy=False)
    745     else:
--> 746         array = np.asarray(array, order=order, dtype=dtype)
    747     except ComplexWarning as complex_warning:
    748         raise ValueError(

~\anaconda3\lib\site-packages\pandas\core\generic.py in __array__(self, dtype)
    2062
    2063     def __array__(self, dtype: npt.DTypeLike | None = None) -> np.ndarray:

```

```
-> 2064         return np.asarray(self._values, dtype=dtype)
    2065
    2066     def __array_wrap__()
```

**ValueError:** could not convert string to float: 'Low'

```
In [134]: 1 # we have to convert strings to numbers before our ML can work  
          2 df_tmp.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 412698 entries, 205615 to 409203
Data columns (total 57 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   SalesID                                   412698 non-null  int64
1   SalePrice                                412698 non-null  float64
2   MachineID                                412698 non-null  int64
3   ModelID                                  412698 non-null  int64
4   datasouce                                412698 non-null  int64
5   auctioneerID                             392562 non-null  float64
6   YearMade                                 412698 non-null  int64
7   MachineHoursCurrentMeter                 147504 non-null  float64
8   UsageBand                                73670 non-null   object
9   fiModelDesc                              412698 non-null  object
10  fiBaseModel                              412698 non-null  object
11  fiSecondaryDesc                          271971 non-null  object
12  fiModelSeries                           58667 non-null   object
13  fiModelDescriptor                       74816 non-null   object
14  ProductSize                             196093 non-null  object
15  fiProductClassDesc                      412698 non-null  object
16  state                                    412698 non-null  object
17  ProductGroup                             412698 non-null  object
18  ProductGroupDesc                        412698 non-null  object
19  Drive_System                            107087 non-null  object
20  Enclosure                               412364 non-null  object
21  Forks                                   197715 non-null  object
22  Pad_Type                                81096 non-null   object
23  Ride_Control                            152728 non-null  object
24  Stick                                   81096 non-null   object
25  Transmission                            188007 non-null  object
26  Turbocharged                            81096 non-null   object
27  Blade_Extension                         25983 non-null   object
28  Blade_Width                             25983 non-null   object
29  Enclosure_Type                         25983 non-null   object
30  Engine_Horsepower                       25983 non-null   object
31  Hydraulics                              330133 non-null  object
32  Pushblock                              25983 non-null   object
33  Ripper                                  106945 non-null  object
34  Scarifier                               25994 non-null   object
35  Tip_Control                             25983 non-null   object
36  Tire_Size                               97638 non-null   object
37  Coupler                                 220679 non-null  object

```

38	Coupler_System	44974	non-null	object
39	Grouser_Tracks	44875	non-null	object
40	Hydraulics_Flow	44875	non-null	object
41	Track_Type	102193	non-null	object
42	Undercarriage_Pad_Width	102916	non-null	object
43	Stick_Length	102261	non-null	object
44	Thumb	102332	non-null	object
45	Pattern_Changer	102261	non-null	object
46	Grouser_Type	102193	non-null	object
47	Backhoe_Mounting	80712	non-null	object
48	Blade_Type	81875	non-null	object
49	Travel_Controls	81877	non-null	object
50	Differential_Type	71564	non-null	object
51	Steering_Controls	71522	non-null	object
52	saleYear	412698	non-null	int64
53	saleMonth	412698	non-null	int64
54	saleDay	412698	non-null	int64
55	saleDayOfWeek	412698	non-null	int64
56	saleDayOfYear	412698	non-null	int64

dtypes: float64(3), int64(10), object(44)  
memory usage: 182.6+ MB

In [135]:

```
1 # A lot of missing data  
2 df_tmp.isna().sum()
```

```
Out[135]: SalesID          0
          SalePrice        0
          MachineID        0
          ModelID          0
          datasource        0
          auctioneerID     20136
          YearMade         0
          MachineHoursCurrentMeter 265194
          UsageBand        339028
          fiModelDesc       0
          fiBaseModel       0
          fiSecondaryDesc   140727
          fiModelSeries     354031
          fiModelDescriptor 337882
          ProductSize       216605
          fiProductClassDesc 0
          state             0
          ProductGroup      0
          ProductGroupDesc  0
          Drive_System      305611
          Enclosure         334
          Forks             214983
          Pad_Type          331602
          Ride_Control      259970
          Stick             331602
          Transmission      224691
          Turbocharged      331602
          Blade_Extension   386715
          Blade_Width       386715
          Enclosure_Type    386715
          Engine_Horsepower 386715
          Hydraulics        82565
          Pushblock         386715
          Ripper            305753
          Scarifier         386704
          Tip_Control       386715
          Tire_Size         315060
          Coupler           192019
          Coupler_System    367724
          Grouser_Tracks    367823
          Hydraulics_Flow   367823
          Track_Type        310505
          Undercarriage_Pad_Width 309782
```



Stick_Length	310437
Thumb	310366
Pattern_Changer	310437
Grouser_Type	310505
Backhoe_Mounting	331986
Blade_Type	330823
Travel_Controls	330821
Differential_Type	341134
Steering_Controls	341176
saleYear	0
saleMonth	0
saleDay	0
saleDayOfWeek	0
saleDayOfYear	0
dtype: int64	

## Convert string to categories

One way we can turn all of our data into numbers is by converting them into pandas categories.

In [136]: 1 df\_tmp.head().T

Out[136]:

	205615	274835	141296	212552	62755
<b>SalesID</b>	1646770	1821514	1505138	1671174	1329056
<b>SalePrice</b>	9500.0	14000.0	50000.0	16000.0	22000.0
<b>MachineID</b>	1126363	1194089	1473654	1327630	1336053
<b>ModelID</b>	8434	10150	4139	8591	4089
<b>datasource</b>	132	132	132	132	132
<b>auctioneerID</b>	18.0	99.0	99.0	99.0	99.0
<b>YearMade</b>	1974	1980	1978	1980	1984
<b>MachineHoursCurrentMeter</b>	NaN	NaN	NaN	NaN	NaN
<b>UsageBand</b>	NaN	NaN	NaN	NaN	NaN
<b>fiModelDesc</b>	TD20	A66	D7G	A62	D3B
<b>fiBaseModel</b>	TD20	A66	D7	A62	D3

In [137]: 1 pd.api.types.is\_string\_dtype(df\_tmp['UsageBand'])

Out[137]: True

```
In [138]: 1 # Find the columns which contain strings
          2 for label, content in df_tmp.items():
          3     if pd.api.types.is_string_dtype(content):
          4         print(label)
```

UsageBand  
fiModelDesc  
fiBaseModel  
fiSecondaryDesc  
fiModelSeries  
fiModelDescriptor  
ProductSize  
fiProductClassDesc  
state  
ProductGroup  
ProductGroupDesc  
Drive\_System  
Enclosure  
Forks  
Pad\_Type  
Ride\_Control  
Stick  
Transmission  
Turbocharged  
Blade\_Extension  
Blade\_Width  
Enclosure\_Type  
Engine\_Horsepower  
Hydraulics  
Pushblock  
Ripper  
Scarifier  
Tip\_Control  
Tire\_Size  
Coupler  
Coupler\_System  
Grouser\_Tracks  
Hydraulics\_Flow  
Track\_Type  
Undercarriage\_Pad\_Width  
Stick\_Length  
Thumb  
Pattern\_Changer  
Grouser\_Type  
Backhoe\_Mounting  
Blade\_Type  
Travel\_Controls

Differential\_Type  
Steering\_Controls

```
In [139]: 1 # if you're wondering what df.items() does, here's an example
2 random_dict = {'key1': 'hello',
3               'key2': 'world!'}
4 for key, value in random_dict.items():
5     print(f'this is a key: {key}',
6           f'this is a value: {value}' )
```

this is a key: key1 this is a value: hello  
this is a key: key2 this is a value: world!

```
In [140]: 1 # This will turn all of the string value into category values
2 for label, content in df_tmp.items():
3     if pd.api.types.is_string_dtype(content):
4         df_tmp[label] = content.astype('category').cat.as_ordered()
```

In [141]:

1	<code>df_tmp.info()</code>
---	----------------------------

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 412698 entries, 205615 to 409203
```

```
Data columns (total 57 columns):
```

#	Column	Non-Null Count	Dtype
0	SalesID	412698 non-null	int64
1	SalePrice	412698 non-null	float64
2	MachineID	412698 non-null	int64
3	ModelID	412698 non-null	int64
4	datasource	412698 non-null	int64
5	auctioneerID	392562 non-null	float64
6	YearMade	412698 non-null	int64
7	MachineHoursCurrentMeter	147504 non-null	float64
8	UsageBand	73670 non-null	category
9	fiModelDesc	412698 non-null	category
10	fiBaseModel	412698 non-null	category
11	fiSecondaryDesc	271971 non-null	category
12	fiModelSeries	58667 non-null	category
13	fiModelDescriptor	74816 non-null	category
14	ProductSize	196093 non-null	category
15	fiProductClassDesc	412698 non-null	category
16	state	412698 non-null	category
17	ProductGroup	412698 non-null	category
18	ProductGroupDesc	412698 non-null	category
19	Drive_System	107087 non-null	category
20	Enclosure	412364 non-null	category
21	Forks	197715 non-null	category
22	Pad_Type	81096 non-null	category
23	Ride_Control	152728 non-null	category
24	Stick	81096 non-null	category
25	Transmission	188007 non-null	category
26	Turbocharged	81096 non-null	category
27	Blade_Extension	25983 non-null	category
28	Blade_Width	25983 non-null	category
29	Enclosure_Type	25983 non-null	category
30	Engine_Horsepower	25983 non-null	category
31	Hydraulics	330133 non-null	category
32	Pushblock	25983 non-null	category
33	Ripper	106945 non-null	category
34	Scarifier	25994 non-null	category
35	Tip_Control	25983 non-null	category
36	Tire_Size	97638 non-null	category
37	Coupler	220679 non-null	category

```

38 Coupler_System          44974 non-null  category
39 Grouser_Tracks          44875 non-null  category
40 Hydraulics_Flow         44875 non-null  category
41 Track_Type              102193 non-null category
42 Undercarriage_Pad_Width 102916 non-null  category
43 Stick_Length            102261 non-null  category
44 Thumb                   102332 non-null  category
45 Pattern_Changer         102261 non-null  category
46 Grouser_Type            102193 non-null  category
47 Backhoe_Mounting        80712 non-null  category
48 Blade_Type              81875 non-null  category
49 Travel_Controls         81877 non-null  category
50 Differential_Type        71564 non-null  category
51 Steering_Controls        71522 non-null  category
52 saleYear                 412698 non-null int64
53 saleMonth                412698 non-null int64
54 saleDay                  412698 non-null int64
55 saleDayOfWeek            412698 non-null int64
56 saleDayOfYear            412698 non-null int64
dtypes: category(44), float64(3), int64(10)
memory usage: 63.2 MB

```

In [142]: 1 df\_tmp.state.cat.categories

```

Out[142]: Index(['Alabama', 'Alaska', 'Arizona', 'Arkansas', 'California', 'Colorado',
'Connecticut', 'Delaware', 'Florida', 'Georgia', 'Hawaii', 'Idaho',
'Illinois', 'Indiana', 'Iowa', 'Kansas', 'Kentucky', 'Louisiana',
'Maine', 'Maryland', 'Massachusetts', 'Michigan', 'Minnesota',
'Mississippi', 'Missouri', 'Montana', 'Nebraska', 'Nevada',
'New Hampshire', 'New Jersey', 'New Mexico', 'New York',
'North Carolina', 'North Dakota', 'Ohio', 'Oklahoma', 'Oregon',
'Pennsylvania', 'Puerto Rico', 'Rhode Island', 'South Carolina',
'South Dakota', 'Tennessee', 'Texas', 'Unspecified', 'Utah', 'Vermont',
'Virginia', 'Washington', 'Washington DC', 'West Virginia', 'Wisconsin',
'Wyoming'],
dtype='object')

```



```
In [143]: 1 df_tmp.state.cat.codes
```

```
Out[143]: 205615    43
          274835     8
          141296     8
          212552     8
          62755     8
          ..
          410879     4
          412476     4
          411927     4
          407124     4
          409203     4
Length: 412698, dtype: int8
```

Thanks to pandas categories, we now have a way to access all of our data in the form of numbers.

But we still have a bunch of missing data...

```
In [144]: 1 df_tmp.isnull().sum()/len(df_tmp)
```

```
Out[144]: SalesID          0.000000
          SalePrice        0.000000
          MachineID        0.000000
          ModelID          0.000000
          datasource        0.000000
          auctioneerID      0.048791
          YearMade          0.000000
          MachineHoursCurrentMeter 0.642586
          UsageBand         0.821492
          fiModelDesc       0.000000
          fiBaseModel       0.000000
          fiSecondaryDesc   0.340993
          fiModelSeries     0.857845
          fiModelDescriptor 0.818715
          ProductSize       0.524851
          fiProductClassDesc 0.000000
          state             0.000000
          ProductGroup      0.000000
          ProductGroupDesc  0.000000
          Drive_System      0.740520
          Enclosure         0.000809
          Forks             0.520921
          Pad_Type          0.803498
          Ride_Control      0.629928
          Stick             0.803498
          Transmission      0.544444
          Turbocharged      0.803498
          Blade_Extension   0.937041
          Blade_Width       0.937041
          Enclosure_Type    0.937041
          Engine_Horsepower 0.937041
          Hydraulics        0.200062
          Pushblock         0.937041
          Ripper            0.740864
          Scarifier         0.937014
          Tip_Control       0.937041
          Tire_Size         0.763415
          Coupler           0.465277
          Coupler_System    0.891024
          Grouser_Tracks    0.891264
          Hydraulics_Flow   0.891264
          Track_Type        0.752378
          Undercarriage_Pad_Width 0.750626
```

Stick_Length	0.752213
Thumb	0.752041
Pattern_Changer	0.752213
Grouser_Type	0.752378
Backhoe_Mounting	0.804428
Blade_Type	0.801610
Travel_Controls	0.801606
Differential_Type	0.826595
Steering_Controls	0.826697
saleYear	0.000000
saleMonth	0.000000
saleDay	0.000000
saleDayOfWeek	0.000000
saleDayOfYear	0.000000
dtype:	float64

## save preprocessed data

```
In [145]: 1 # eExport current tmp dataframe
          2 df_tmp.to_csv('train_tmp.csv',
          3                  index=False)
```

In [146]:

```

1 # Import preprocessed data
2
3 df_tmp = pd.read_csv('train_tmp.csv',
4                       low_memory=False)
5 df_tmp.head().T

```

Out[146]:

	0	1	2	3	4
<b>SalesID</b>	1646770	1821514	1505138	1671174	1329056
<b>SalePrice</b>	9500.0	14000.0	50000.0	16000.0	22000.0
<b>MachineID</b>	1126363	1194089	1473654	1327630	1336053
<b>ModelID</b>	8434	10150	4139	8591	4089
<b>datasource</b>	132	132	132	132	132
<b>auctioneerID</b>	18.0	99.0	99.0	99.0	99.0
<b>YearMade</b>	1974	1980	1978	1980	1984
<b>MachineHoursCurrentMeter</b>	NaN	NaN	NaN	NaN	NaN
<b>UsageBand</b>	NaN	NaN	NaN	NaN	NaN
<b>fiModelDesc</b>	TD20	A66	D7G	A62	D3B
<b>fiBaseModel</b>	TD20	A66	D7	A62	D3

In [147]:

1	<code>df_tmp.isna().sum()</code>
---	----------------------------------

```
Out[147]: SalesID          0
          SalePrice        0
          MachineID        0
          ModelID          0
          datasource        0
          auctioneerID     20136
          YearMade         0
          MachineHoursCurrentMeter 265194
          UsageBand        339028
          fiModelDesc       0
          fiBaseModel       0
          fiSecondaryDesc   140727
          fiModelSeries     354031
          fiModelDescriptor 337882
          ProductSize       216605
          fiProductClassDesc 0
          state             0
          ProductGroup      0
          ProductGroupDesc  0
          Drive_System      305611
          Enclosure         334
          Forks             214983
          Pad_Type          331602
          Ride_Control      259970
          Stick             331602
          Transmission      224691
          Turbocharged      331602
          Blade_Extension   386715
          Blade_Width       386715
          Enclosure_Type    386715
          Engine_Horsepower 386715
          Hydraulics        82565
          Pushblock         386715
          Ripper            305753
          Scarifier         386704
          Tip_Control       386715
          Tire_Size         315060
          Coupler           192019
          Coupler_System    367724
          Grouser_Tracks    367823
          Hydraulics_Flow   367823
          Track_Type        310505
          Undercarriage_Pad_Width 309782
```

```
Stick_Length      310437
Thumb             310366
Pattern_Changer   310437
Grouser_Type      310505
Backhoe_Mounting  331986
Blade_Type        330823
Travel_Controls   330821
Differential_Type  341134
Steering_Controls 341176
saleYear          0
saleMonth         0
saleDay           0
saleDayOfWeek     0
saleDayOfYear     0
dtype: int64
```

## Fill missing values

### Fill numerical missing values first

```
In [148]: 1 for label, content in df_tmp.items():
          2     if pd.api.types.is_numeric_dtype(content):
          3         print(label)
```

```
SalesID
SalePrice
MachineID
ModelID
datasource
auctioneerID
YearMade
MachineHoursCurrentMeter
saleYear
saleMonth
saleDay
saleDayOfWeek
saleDayOfYear
```



In [149]: 1 df\_tmp.ModelID

Out[149]: 0 8434  
1 10150  
2 4139  
3 8591  
4 4089  
...  
412693 5266  
412694 19330  
412695 17244  
412696 3357  
412697 4701

Name: ModelID, Length: 412698, dtype: int64

```
In [150]: 1 # check for which numeric clumns have null values
2 for label, content in df_tmp.items():
3     if pd.api.types.is_numeric_dtype(content):
4         if pd.isnull(content).sum():
5             print(label)
6
```

auctioneerID  
MachineHoursCurrentMeter

```
In [151]: 1 # Fill numeric rows with the median
2 for label, content in df_tmp.items():
3     if pd.api.types.is_numeric_dtype(content):
4         if pd.isnull(content).sum():
5             # Add a binary column which tells us if the data was missing or not
6             df_tmp[label+'_is_missing'] = pd.isnull(content)
7             # Fill missing numeric values with median
8             df_tmp[label] = content.fillna(content.median())
9
```

```
In [152]: 1 # Demonstrate how median is more robust than mean
          2 hundreds = np.full((1000), 100)
          3 hundreds_billion = np.append(hundreds, 10000000000)
          4 np.mean(hundreds), np.mean(hundreds_billion), np.median(hundreds), np.median(hundreds_billion)
```

```
Out[152]: (100.0, 999100.8991008991, 100.0, 100.0)
```

```
In [153]: 1 for label, content in df_tmp.items():
          2     if pd.api.types.is_numeric_dtype(content):
          3         if pd.isnull(content).sum():
          4             print(label)
```

```
In [154]: 1 # cheeck to see how many examples were missing
          2 df_tmp.auctioneerID_is_missing.value_counts(), df_tmp.MachineHoursCurrentMeter_is_missing.value_count
```

```
Out[154]: (False      392562
          True        20136
          Name: auctioneerID_is_missing, dtype: int64,
          True        265194
          False       147504
          Name: MachineHoursCurrentMeter_is_missing, dtype: int64)
```

In [155]:

1	<code>df_tmp.isna().sum()</code>
---	----------------------------------

```

Out[155]: SalesID                0
          SalePrice              0
          MachineID             0
          ModelID               0
          datasource            0
          auctioneerID         0
          YearMade              0
          MachineHoursCurrentMeter 0
          UsageBand            339028
          fiModelDesc           0
          fiBaseModel           0
          fiSecondaryDesc       140727
          fiModelSeries         354031
          fiModelDescriptor     337882
          ProductSize           216605
          fiProductClassDesc    0
          state                 0
          ProductGroup          0
          ProductGroupDesc      0
          Drive_System          305611
          Enclosure             334
          Forks                 214983
          Pad_Type              331602
          Ride_Control          259970
          Stick                 331602
          Transmission          224691
          Turbocharged          331602
          Blade_Extension       386715
          Blade_Width           386715
          Enclosure_Type       386715
          Engine_Horsepower     386715
          Hydraulics            82565
          Pushblock             386715
          Ripper                305753
          Scarifier             386704
          Tip_Control           386715
          Tire_Size             315060
          Coupler               192019
          Coupler_System        367724
          Grouser_Tracks        367823
          Hydraulics_Flow       367823
          Track_Type            310505
          Undercarriage_Pad_Width 309782

```

Stick_Length	310437
Thumb	310366
Pattern_Changer	310437
Grouser_Type	310505
Backhoe_Mounting	331986
Blade_Type	330823
Travel_Controls	330821
Differential_Type	341134
Steering_Controls	341176
saleYear	0
saleMonth	0
saleDay	0
saleDayOfWeek	0
saleDayOfYear	0
auctioneerID_is_missing	0
MachineHoursCurrentMeter_is_missing	0

dtype: int64

## Filling and turning categorical variables into numbers

```
In [156]: 1 # check for columns which aren't numeric
          2 for label, content in df_tmp.items():
          3     if not pd.api.types.is_numeric_dtype(content):
          4         print(label)
```

UsageBand  
fiModelDesc  
fiBaseModel  
fiSecondaryDesc  
fiModelSeries  
fiModelDescriptor  
ProductSize  
fiProductClassDesc  
state  
ProductGroup  
ProductGroupDesc  
Drive\_System  
Enclosure  
Forks  
Pad\_Type  
Ride\_Control  
Stick  
Transmission  
Turbocharged  
Blade\_Extension  
Blade\_Width  
Enclosure\_Type  
Engine\_Horsepower  
Hydraulics  
Pushblock  
Ripper  
Scarifier  
Tip\_Control  
Tire\_Size  
Coupler  
Coupler\_System  
Grouser\_Tracks  
Hydraulics\_Flow  
Track\_Type  
Undercarriage\_Pad\_Width  
Stick\_Length  
Thumb  
Pattern\_Changer  
Grouser\_Type  
Backhoe\_Mounting  
Blade\_Type  
Travel\_Controls

Differential\_Type  
Steering\_Controls

```
In [157]: 1 # Turn categorical variables into numbers and fill missing
          2 for label, content in df_tmp.items():
          3     if not pd.api.types.is_numeric_dtype(content):
          4         # Add binary column to indicate whether sample has missing value
          5         df_tmp[label+'_is_missing'] = pd.isnull(content)
          6         # Turn categories into numbers and add +1
          7         df_tmp[label] = pd.Categorical(content).codes + 1
```

```
In [158]: 1 pd.Categorical(df_tmp['state']).codes
```

```
Out[158]: array([43,  8,  8, ...,  4,  4,  4], dtype=int8)
```

```
In [159]: 1 df_tmp.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 412698 entries, 0 to 412697
Columns: 103 entries, SalesID to Steering_Controls_is_missing
dtypes: bool(46), float64(3), int16(4), int64(10), int8(40)
memory usage: 77.9 MB
```



In [160]: 1 df\_tmp.head().T

Out[160]:

	0	1	2	3	4
<b>SalesID</b>	1646770	1821514	1505138	1671174	1329056
<b>SalePrice</b>	9500.0	14000.0	50000.0	16000.0	22000.0
<b>MachineID</b>	1126363	1194089	1473654	1327630	1336053
<b>ModelID</b>	8434	10150	4139	8591	4089
<b>datasource</b>	132	132	132	132	132
...	...	...	...	...	...
<b>Backhoe_Mounting_is_missing</b>	False	True	False	True	False
<b>Blade_Type_is_missing</b>	False	True	False	True	False
<b>Travel_Controls_is_missing</b>	False	True	False	True	False
<b>Differential_Type_is_missing</b>	True	False	True	False	True
<b>Steering_Controls_is_missing</b>	True	False	True	False	True

103 rows × 5 columns

In [161]: 1 df\_tmp.isna().sum()

Out[161]:

SalesID	0
SalePrice	0
MachineID	0
ModelID	0
datasource	0
..	
Backhoe_Mounting_is_missing	0
Blade_Type_is_missing	0
Travel_Controls_is_missing	0
Differential_Type_is_missing	0
Steering_Controls_is_missing	0
Length: 103, dtype: int64	

```
In [162]: 1 pd.Categorical(df_tmp['UsageBand']).codes
```

```
Out[162]: array([0, 0, 0, ..., 0, 0, 0], dtype=int8)
```

```
In [163]: 1 pd.Categorical(df_tmp['UsageBand']).codes + 1
```

```
Out[163]: array([1, 1, 1, ..., 1, 1, 1], dtype=int8)
```


Now that all of data is numeric as well as our dataframe has no missing values, we should be able to build a machine learning model

```
In [164]: 1 df_tmp.head()
```

```
Out[164]:
```

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter	UsageBand	fiModelDesc	
0	1646770	9500.0	1126363	8434	132	18.0	1974	0.0	0	4593	.
1	1821514	14000.0	1194089	10150	132	99.0	1980	0.0	0	1820	.
2	1505138	50000.0	1473654	4139	132	99.0	1978	0.0	0	2348	.
3	1671174	16000.0	1327630	8591	132	99.0	1980	0.0	0	1819	.
4	1329056	22000.0	1336053	4089	132	99.0	1984	0.0	0	2119	.

5 rows × 103 columns



```
In [165]: 1 len(df_tmp)
```

```
Out[165]: 412698
```

```
In [166]: 1 %%time
          2 # Instantiate model
          3 model = RandomForestRegressor(n_jobs=-1,
          4                               random_state=42)
          5
          6 # Fit the model
          7 model.fit(df_tmp.drop('SalePrice', axis = 1), df_tmp['SalePrice'])
```

Wall time: 10min 25s

Out[166]: RandomForestRegressor(n\_jobs=-1, random\_state=42)

```
In [167]: 1 model.score(df_tmp.drop('SalePrice', axis = 1), df_tmp['SalePrice'])
```

Out[167]: 0.9875468079970562

**Question:** Why doesn't the above metric hold water? (why isn't metric reliable)

## Splitting data into training/validation sets

```
In [168]: 1 df_tmp.saleYear
```

Out[168]:

0	1989
1	1989
2	1989
3	1989
4	1989
	...
412693	2012
412694	2012
412695	2012
412696	2012
412697	2012

Name: saleYear, Length: 412698, dtype: int64

```
In [169]: 1 df_tmp.saleYear.value_counts()
```

```
Out[169]: 2009    43849
          2008    39767
          2011    35197
          2010    33390
          2007    32208
          2006    21685
          2005    20463
          2004    19879
          2001    17594
          2000    17415
          2002    17246
          2003    15254
          1998    13046
          1999    12793
          2012    11573
          1997     9785
          1996     8829
          1995     8530
          1994     7929
          1993     6303
          1992     5519
          1991     5109
          1989     4806
          1990     4529
          Name: saleYear, dtype: int64
```

```
In [170]: 1 # split data into training and validation
          2 df_val = df_tmp[df_tmp.saleYear == 2012]
          3 df_train = df_tmp[df_tmp.saleYear != 2012]
          4
          5 len(df_val), len(df_train)
```

```
Out[170]: (11573, 401125)
```

```
In [171]: 1 # split data into x and y
          2 x_train, y_train = df_train.drop('SalePrice', axis = 1), df_train.SalePrice
          3 x_valid, y_valid = df_val.drop('SalePrice', axis = 1), df_val.SalePrice
          4
          5 x_train.shape, y_train.shape, x_valid.shape, y_valid.shape
```

```
Out[171]: ((401125, 102), (401125,), (11573, 102), (11573,))
```

```
In [172]: 1 y_train
```


```
Out[172]: 0          9500.0
          1         14000.0
          2         50000.0
          3         16000.0
          4         22000.0
          ...
          401120        29000.0
          401121        11000.0
          401122        11000.0
          401123        18000.0
          401124        13500.0
          Name: SalePrice, Length: 401125, dtype: float64
```

In [173]: 1 x\_train

Out[173]:

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter	UsageBand	fiModelDesc	fiBase
0	1646770	1126363	8434	132	18.0	1974	0.0	0	4593	
1	1821514	1194089	10150	132	99.0	1980	0.0	0	1820	
2	1505138	1473654	4139	132	99.0	1978	0.0	0	2348	
3	1671174	1327630	8591	132	99.0	1980	0.0	0	1819	
4	1329056	1336053	4089	132	99.0	1984	0.0	0	2119	
...	...	...	...	...	...	...	...	...	...	...
401120	6260687	1074871	4331	149	2.0	1000	0.0	0	3137	
401121	6312170	1812622	9580	149	2.0	2005	0.0	0	4514	
401122	6312727	1811599	9580	149	2.0	2005	0.0	0	4514	
401123	6315051	1858173	17432	149	2.0	2004	0.0	0	3389	
401124	6260878	1799594	4102	149	2.0	1000	0.0	0	2161	

401125 rows × 102 columns




In [174]:

1 x\_valid

Out[174]:

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter	UsageBand	fiModelDesc	fiBase
<b>401125</b>	4449186	2318824	26964	173	99.0	1996	0.0	0	2269	
<b>401126</b>	1222855	531393	23926	121	3.0	1000	8145.0	2	85	
<b>401127</b>	6258613	1810917	13260	149	99.0	2000	24.0	2	1115	
<b>401128</b>	6282680	1543404	1830	149	99.0	2004	4373.0	3	64	
<b>401129</b>	6282759	1863077	11390	149	99.0	2006	3467.0	3	139	
...	...	...	...	...	...	...	...	...	...	...
<b>412693</b>	6302984	1915521	5266	149	99.0	2001	0.0	0	2101	
<b>412694</b>	6324811	1919104	19330	149	99.0	2004	0.0	0	240	
<b>412695</b>	6313029	1918416	17244	149	99.0	2004	0.0	0	627	
<b>412696</b>	6266251	509560	3357	149	99.0	1993	0.0	0	83	
<b>412697</b>	6283635	1869284	4701	149	99.0	1000	0.0	0	989	

11573 rows × 102 columns



In [175]:

1 y\_valid

Out[175]:

401125	46173.2
401126	66000.0
401127	26800.0
401128	42100.0
401129	62100.0

...

412693	16000.0
412694	6000.0
412695	16000.0
412696	55000.0
412697	34000.0

Name: SalePrice, Length: 11573, dtype: float64

## Building an evaluation function

```
In [ ]: 1 # create evaluation function (the competition uses RMSLE)
2 from sklearn.metrics import mean_squared_log_error, mean_absolute_error, r2_score
3
4 def rmsle(y_test, y_preds):
5     '''
6     calculate root mean squared log error between predictions and true labels'''
7     return np.sqrt(mean_squared_log_error(y_test, y_preds))
8
9
10 # create function to evaluate model on a few different levels
11 def show_scores(model):
12     train_preds = model.predict(x_train)
13     val_preds = model.predict(x_valid)
14     scores = {'Training MAE': mean_absolute_error(y_train, train_preds),
15              'Valid MAE': mean_absolute_error(y_valid, val_preds),
16              'Training RMSLE': rmsle(y_train, train_preds),
17              'Valid RMSLE': rmsle(y_valid, val_preds),
18              'Training R^2': r2_score(y_train, train_preds),
19              'Valid R^2': r2_score(y_valid, val_preds)}
20     return scores
```

## Testing our model on a subset (to tune the hyperparameters)

```
In [ ]: 1 # # This takes far too long... for experimenting
2 # %%time
3 # model = RandomForestRegressor(n_jobs= -1,
4 #                               random_state=42)
5 # # model.fit(x_train, y_train)
```

```
In [ ]: 1 len(x_train)
```



```
In [ ]: 1 # Change max_samples value
        2 model = RandomForestRegressor(n_jobs=-1,
        3                               random_state=42,
        4                               max_samples=10000)
```

```
In [ ]: 1 %%time
        2 # cutting down on the max number of samples each estimator can see improves training time
        3 model.fit(x_train, y_train)
```

```
In [ ]: 1 show_scores(model)
```

## Hyperparameter tuning with RandomizedSearchCV

```

In [182]: 1 %%time
2 from sklearn.model_selection import RandomizedSearchCV
3
4 # Different RandomForestRegressor hyperparameters
5
6 rf_grid = {'n_estimators': np.arange(10, 100, 10),
7            'max_depth': [None, 3, 5, 10],
8            'min_samples_split': np.arange(2, 20, 2),
9            'min_samples_leaf': np.arange(1, 20, 2),
10           'max_features': [0.5, 1, 'sqrt', 'auto'],
11           'max_samples': [10000]}
12
13 # instantiate RandomizedSearchCV model
14 rs_model = RandomizedSearchCV(RandomForestRegressor(n_jobs = -1,
15                                                    random_state=42),
16                               param_distributions=rf_grid,
17                               n_iter = 2,
18                               cv=5,
19                               verbose=True)
20 # Fit the RandomizedSearchCV
21 rs_model.fit(x_train, y_train)

```

Fitting 5 folds for each of 2 candidates, totalling 10 fits

Wall time: 3min 24s

Parser : 186 ms

```

Out[182]: RandomizedSearchCV(cv=5,
                             estimator=RandomForestRegressor(n_jobs=-1, random_state=42),
                             n_iter=2,
                             param_distributions={'max_depth': [None, 3, 5, 10],
                                                  'max_features': [0.5, 1, 'sqrt',
                                                                'auto'],
                                                  'max_samples': [10000],
                                                  'min_samples_leaf': array([ 1,  3,  5,  7,  9, 11, 13, 15, 17, 1
9]),
                                                  'min_samples_split': array([ 2,  4,  6,  8, 10, 12, 14, 16, 1
8]),
                                                  'n_estimators': array([10, 20, 30, 40, 50, 60, 70, 80, 90])},
                             verbose=True)

```

```
In [183]: 1 # find the best model hyperparameters
          2 rs_model.best_params_
```

```
Out[183]: {'n_estimators': 80,
           'min_samples_split': 10,
           'min_samples_leaf': 3,
           'max_samples': 10000,
           'max_features': 'sqrt',
           'max_depth': 5}
```

```
In [184]: 1 # evaluate the RandomizedSearchCV
          2 show_scores(rs_model)
```

```
Out[184]: {'Training MAE': 11730.789081883375,
           'Valid MAE': 13587.261844750172,
           'Training RMSLE': 0.5045615136209148,
           'Valid RMSLE': 0.5166094651411732,
           'Training R^2': 0.488496697482226,
           'Valid R^2': 0.4867697712132065}
```

## Train a model with the best hyperparameters

**Note** These were found after 100 iterations of RandomizedSaerchCV

```
In [185]: 1 %%time
2
3 # most ideals hyperparameters
4 ideal_model = RandomForestRegressor(n_estimators=40,
5                                     min_samples_leaf=1,
6                                     min_samples_split=14,
7                                     max_features=0.5,
8                                     n_jobs=-1,
9                                     max_samples=None,
10                                    random_state=42)
11
12 # fit the ideal model
13 ideal_model.fit(x_train, y_train)
```

Wall time: 1min 58s

```
Out[185]: RandomForestRegressor(max_features=0.5, min_samples_split=14, n_estimators=40,
                                n_jobs=-1, random_state=42)
```

```
In [186]: 1 # scores for ideal_model (trained on all the data)
2 show_scores(ideal_model)
```

```
Out[186]: {'Training MAE': 2953.8161137163484,
            'Valid MAE': 5951.247761444453,
            'Training RMSLE': 0.14469006962371858,
            'Valid RMSLE': 0.2452416398953833,
            'Training R^2': 0.9588145522577225,
            'Valid R^2': 0.8818019502450094}
```

```
In [187]: 1 # scores on rs_model (only trained on ~10,000 examples)
2 show_scores(rs_model)
```

```
Out[187]: {'Training MAE': 11730.789081883375,
            'Valid MAE': 13587.261844750172,
            'Training RMSLE': 0.5045615136209148,
            'Valid RMSLE': 0.5166094651411733,
            'Training R^2': 0.488496697482226,
            'Valid R^2': 0.4867697712132065}
```


## Make predictions on test data

```
In [221]: 1 # import the test data
          2 df_test = pd.read_csv('Test.csv',
          3                               low_memory=False,
          4                               parse_dates=['saledate'])
          5 df_test.head()
```

Out[221]:

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter	UsageBand	saledate	fiModelDesc	...
0	1227829	1006309	3168	121	3	1999	3688.0	Low	2012-05-03	580G	...
1	1227844	1022817	7271	121	3	1000	28555.0	High	2012-05-10	936	...
2	1227847	1031560	22805	121	3	2004	6038.0	Medium	2012-05-10	EC210BLC	...
3	1227848	56204	1269	121	3	2006	8940.0	High	2012-05-10	330CL	...
4	1227863	1053887	22312	121	3	2005	2286.0	Low	2012-05-10	650K	...

5 rows × 52 columns



In [222]:

```
1 # Make predictions on the test dataset
2 test_preds = ideal_model.predict(df_test)
```

C:\Users\USER\anaconda3\lib\site-packages\sklearn\base.py:493: FutureWarning: The feature names should match those that were passed during fit. Starting version 1.2, an error will be raised.

Feature names unseen at fit time:

- saledate

Feature names seen at fit time, yet now missing:

- Backhoe\_Mounting\_is\_missing
- Blade\_Extension\_is\_missing
- Blade\_Type\_is\_missing
- Blade\_Width\_is\_missing
- Coupler\_System\_is\_missing
- ...

warnings.warn(message, FutureWarning)

```

-----
ValueError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_6476\295568641.py in <module>
      1 # Make predictions on the test dataset
----> 2 test_preds = ideal_model.predict(df_test)

~\anaconda3\lib\site-packages\sklearn\ensemble\_forest.py in predict(self, X)
    969     check_is_fitted(self)
    970     # Check data
--> 971     X = self._validate_X_predict(X)
    972
    973     # Assign chunk of trees to jobs

~\anaconda3\lib\site-packages\sklearn\ensemble\_forest.py in _validate_X_predict(self, X)
    577     Validate X whenever one tries to predict, apply, predict_proba."""
    578     check_is_fitted(self)
--> 579     X = self._validate_data(X, dtype=DTYPE, accept_sparse="csr", reset=False)
    580     if issparse(X) and (X.indices.dtype != np.intc or X.indptr.dtype != np.intc):
    581         raise ValueError("No support for np.int64 index based sparse matrices")

~\anaconda3\lib\site-packages\sklearn\base.py in _validate_data(self, X, y, reset, validate_separately,
**check_params)
    564         raise ValueError("Validation should be done on X, y or both.")
    565     elif not no_val_X and no_val_y:
--> 566         X = check_array(X, **check_params)
    567         out = X
    568     elif no_val_X and not no_val_y:

~\anaconda3\lib\site-packages\sklearn\utils\validation.py in check_array(array, accept_sparse, accept_large_sparse, dtype, order, copy, force_all_finite, ensure_2d, allow_nd, ensure_min_samples, ensure_min_features, estimator)
    744         array = array.astype(dtype, casting="unsafe", copy=False)
    745     else:
--> 746         array = np.asarray(array, order=order, dtype=dtype)
    747     except ComplexWarning as complex_warning:
    748         raise ValueError(

~\anaconda3\lib\site-packages\pandas\core\generic.py in __array__(self, dtype)
    2062
    2063     def __array__(self, dtype: npt.DTypeLike | None = None) -> np.ndarray:
-> 2064         return np.asarray(self._values, dtype=dtype)
    2065
    2066     def __array_wrap__(

```

**ValueError:** could not convert string to float: 'Low'

**Preprocessing the data (getting the test dataset in the same format as our training dataset)**



In [223]:

1	<code>df_test.isna().sum()</code>
---	-----------------------------------

```

Out[223]: SalesID          0
          MachineID       0
          ModelID         0
          datasource      0
          auctioneerID    0
          YearMade        0
          MachineHoursCurrentMeter 10328
          UsageBand      10623
          saledate        0
          fiModelDesc     0
          fiBaseModel     0
          fiSecondaryDesc 3975
          fiModelSeries   10451
          fiModelDescriptor 9433
          ProductSize     6409
          fiProductClassDesc 0
          state           0
          ProductGroup    0
          ProductGroupDesc 0
          Drive_System    9698
          Enclosure       2
          Forks           6149
          Pad_Type       10349
          Ride_Control    8216
          Stick           10349
          Transmission    7639
          Turbocharged    10349
          Blade_Extension 11806
          Blade_Width     11806
          Enclosure_Type  11806
          Engine_Horsepower 11806
          Hydraulics      2142
          Pushblock       11806
          Ripper          9753
          Scarifier       11806
          Tip_Control     11806
          Tire_Size       9679
          Coupler         4856
          Coupler_System  10391
          Grouser_Tracks  10391
          Hydraulics_Flow 10391
          Track_Type      9063
          Undercarriage_Pad_Width 9059

```

```

Stick_Length          9063
Thumb                 9062
Pattern_Changer       9063
Grouser_Type          9063
Backhoe_Mounting      10406
Blade_Type            10399
Travel_Controls       10399
Differential_Type     10328
Steering_Controls     10328
dtype: int64

```

In [224]: 1 df\_test.columns

```

Out[224]: Index(['SalesID', 'MachineID', 'ModelID', 'datasource', 'auctioneerID',
                'YearMade', 'MachineHoursCurrentMeter', 'UsageBand', 'saledate',
                'fiModelDesc', 'fiBaseModel', 'fiSecondaryDesc', 'fiModelSeries',
                'fiModelDescriptor', 'ProductSize', 'fiProductClassDesc', 'state',
                'ProductGroup', 'ProductGroupDesc', 'Drive_System', 'Enclosure',
                'Forks', 'Pad_Type', 'Ride_Control', 'Stick', 'Transmission',
                'Turbocharged', 'Blade_Extension', 'Blade_Width', 'Enclosure_Type',
                'Engine_Horsepower', 'Hydraulics', 'Pushblock', 'Ripper', 'Scarifier',
                'Tip_Control', 'Tire_Size', 'Coupler', 'Coupler_System',
                'Grouser_Tracks', 'Hydraulics_Flow', 'Track_Type',
                'Undercarriage_Pad_Width', 'Stick_Length', 'Thumb', 'Pattern_Changer',
                'Grouser_Type', 'Backhoe_Mounting', 'Blade_Type', 'Travel_Controls',
                'Differential_Type', 'Steering_Controls'],
                dtype='object')

```

In [225]: 1 x\_train.columns

```

Out[225]: Index(['SalesID', 'MachineID', 'ModelID', 'datasource', 'auctioneerID',
                'YearMade', 'MachineHoursCurrentMeter', 'UsageBand', 'fiModelDesc',
                'fiBaseModel',
                ...,
                'Undercarriage_Pad_Width_is_missing', 'Stick_Length_is_missing',
                'Thumb_is_missing', 'Pattern_Changer_is_missing',
                'Grouser_Type_is_missing', 'Backhoe_Mounting_is_missing',
                'Blade_Type_is_missing', 'Travel_Controls_is_missing',
                'Differential_Type_is_missing', 'Steering_Controls_is_missing'],
                dtype='object', length=102)

```

In [226]:

```
1 def preprocess_data(df):
2     '''
3     performs transformations on df and returns transformed df.
4     '''
5     df['saleYear'] = df.saledate.dt.year
6     df['saleMonth'] = df.saledate.dt.month
7     df['saleDay'] = df.saledate.dt.day
8     df['saleDayOfWeek'] = df.saledate.dt.dayofweek
9     df['saleDayOfYear'] = df.saledate.dt.dayofyear
10
11     df.drop('saledate', axis = 1, inplace=True)
12
13     # Fill the numeric rows with median
14     for label, content in df.items():
15         if pd.api.types.is_numeric_dtype(content):
16             if pd.isnull(content).sum():
17                 # Add a binary column which tells us if the data was missing or not
18                 df[label+'_is_missing'] = pd.isnull(content)
19                 # Fill missing numeric values with median
20                 df[label] = content.fillna(content.median())
21     # This will turn all of the string value into category values
22     for label, content in df.items():
23         if pd.api.types.is_string_dtype(content):
24             df[label] = content.astype('category').cat.as_ordered()
25
26     # Filled categorical missing data and turn categories into numbers
27     if not pd.api.types.is_numeric_dtype(content):
28         df[label+'_is_missing'] = pd.isnull(content)
29         # we add +1 to the category code because pandas encodes missing categories as -1
30         df[label] = pd.Categorical(content).codes + 1
31
32
33     return df
```

In [227]:

```
1 # process test data
2 df_test = preprocess_data(df_test)
3 df_test.head()
```

Out[227]:

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter	UsageBand	fiModelDesc	fiBaseModel
0	1227829	1006309	3168	121	3	1999	3688.0	2	499	180
1	1227844	1022817	7271	121	3	1000	28555.0	1	831	292
2	1227847	1031560	22805	121	3	2004	6038.0	3	1177	404
3	1227848	56204	1269	121	3	2006	8940.0	1	287	113
4	1227863	1053887	22312	121	3	2005	2286.0	2	566	196

5 rows × 101 columns



In [228]:

```
1 # make predictions on updated test data
2 test_preds = ideal_model.predict(df_test)
```

C:\Users\USER\anaconda3\lib\site-packages\sklearn\base.py:493: FutureWarning: The feature names should match those that were passed during fit. Starting version 1.2, an error will be raised.

Feature names seen at fit time, yet now missing:

- auctioneerID\_is\_missing

```
warnings.warn(message, FutureWarning)
```

```

-----
ValueError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_6476\2973351858.py in <module>
      1 # make predictions on updated test data
----> 2 test_preds = ideal_model.predict(df_test)

~\anaconda3\lib\site-packages\sklearn\ensemble\_forest.py in predict(self, X)
    969         check_is_fitted(self)
    970         # Check data
--> 971         X = self._validate_X_predict(X)
    972
    973         # Assign chunk of trees to jobs

~\anaconda3\lib\site-packages\sklearn\ensemble\_forest.py in _validate_X_predict(self, X)
    577         Validate X whenever one tries to predict, apply, predict_proba."""
    578         check_is_fitted(self)
--> 579         X = self._validate_data(X, dtype=DTYPE, accept_sparse="csr", reset=False)
    580         if issparse(X) and (X.indices.dtype != np.intc or X.indptr.dtype != np.intc):
    581             raise ValueError("No support for np.int64 index based sparse matrices")

~\anaconda3\lib\site-packages\sklearn\base.py in _validate_data(self, X, y, reset, validate_separately,
**check_params)
    583
    584         if not no_val_X and check_params.get("ensure_2d", True):
--> 585             self._check_n_features(X, reset=reset)
    586
    587         return out

~\anaconda3\lib\site-packages\sklearn\base.py in _check_n_features(self, X, reset)
    398
    399         if n_features != self.n_features_in_:
--> 400             raise ValueError(
    401                 f"X has {n_features} features, but {self.__class__.__name__} "
    402                 f"is expecting {self.n_features_in_} features as input."

```

**ValueError:** X has 101 features, but RandomForestRegressor is expecting 102 features as input.


In [229]:

```
1 x_train.head()
```

Out[229]:

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter	UsageBand	fiModelDesc	fiBaseModel
0	1646770	1126363	8434	132	18.0	1974	0.0	0	4593	1744
1	1821514	1194089	10150	132	99.0	1980	0.0	0	1820	559
2	1505138	1473654	4139	132	99.0	1978	0.0	0	2348	713
3	1671174	1327630	8591	132	99.0	1980	0.0	0	1819	558
4	1329056	1336053	4089	132	99.0	1984	0.0	0	2119	683

5 rows × 102 columns



In [230]:

```
1 # we can find how the columns differ using sets
2 set(x_train.columns) - set(df_test.columns)
```

Out[230]: {'auctioneerID\_is\_missing'}


In [231]:

```
1 # Manually adjust df-test to have auctioneerID_is_missing
2 df_test['auctioneerID_is_missing'] = False
3 df_test.head()
```

Out[231]:

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter	UsageBand	fiModelDesc	fiBaseModel
0	1227829	1006309	3168	121	3	1999	3688.0	2	499	180
1	1227844	1022817	7271	121	3	1000	28555.0	1	831	292
2	1227847	1031560	22805	121	3	2004	6038.0	3	1177	404
3	1227848	56204	1269	121	3	2006	8940.0	1	287	113
4	1227863	1053887	22312	121	3	2005	2286.0	2	566	196

5 rows × 102 columns



Finally now our test dataframe has the same features as our training dataframe, we can make predictions



```
In [232]: 1 # make predictions of the test data
          2 test_preds = ideal_model.predict(df_test)
```

C:\Users\USER\anaconda3\lib\site-packages\sklearn\base.py:493: FutureWarning: The feature names should match those that were passed during fit. Starting version 1.2, an error will be raised. Feature names must be in the same order as they were in fit.

```
warnings.warn(message, FutureWarning)
```

```
In [233]: 1 test_preds
```

```
Out[233]: array([20614.36780887, 19897.80170658, 44852.21959446, ...,
                14296.98620472, 22164.85757662, 31683.80063427])
```

we've made some predictions but they're not in the same format Kaggle is asking for.

In [237]:

```

1 # Format predictions into the same format kaggle is after
2 df_preds = pd.DataFrame()
3 df_preds['SalesID'] = df_test['SalesID']
4 df_preds['salesPrice'] = test_preds
5 df_preds

```

Out[237]:

	SalesID	salesPrice
0	1227829	20614.367809
1	1227844	19897.801707
2	1227847	44852.219594
3	1227848	68346.325323
4	1227863	39487.349708
...	...	...
12452	6643171	46466.092910
12453	6643173	17500.493352
12454	6643184	14296.986205
12455	6643186	22164.857577
12456	6643196	31683.800634

12457 rows × 2 columns

In [254]:

```

1 # Export prediction data
2 df_preds.to_csv('test_predictions.csv', index=False)

```

## Feature Importance

Feature importance seeks to figure out which different attribute of the data were most importance when it comes to predicting the **target variable** (SalePrice)

```
In [255]: 1 # find feature Importance of our best model
          2 ideal_model.feature_importances_
```

```
Out[255]: array([3.39445533e-02, 1.81148281e-02, 4.09167072e-02, 1.70752171e-03,
                 3.40797459e-03, 2.08200698e-01, 2.95067052e-03, 1.10113725e-03,
                 4.16122668e-02, 4.71911805e-02, 6.23815431e-02, 4.67433955e-03,
                 1.52524442e-02, 1.52517337e-01, 4.72224713e-02, 5.96817956e-03,
                 1.29351899e-03, 2.78088439e-03, 2.37248769e-03, 6.17114453e-02,
                 8.13525488e-04, 3.61873268e-05, 9.19098115e-04, 2.23170993e-04,
                 1.28102678e-03, 2.06519636e-05, 2.01477316e-03, 6.63364759e-03,
                 2.15274492e-03, 2.50178165e-03, 4.63902393e-03, 3.85873985e-03,
                 2.76062667e-03, 1.00782454e-03, 2.47969268e-04, 6.04239818e-03,
                 7.64997072e-04, 1.57100537e-02, 2.29716203e-03, 2.58372272e-03,
                 8.07637426e-04, 9.18548690e-04, 1.35656446e-03, 5.81458569e-04,
                 4.96716928e-04, 3.79552257e-04, 5.31712788e-04, 2.71823509e-03,
                 8.34294376e-04, 3.12136841e-04, 2.14075157e-04, 7.42422919e-02,
                 3.80158492e-03, 5.67641024e-03, 2.87154703e-03, 9.83349904e-03,
                 2.65470837e-04, 1.57946459e-03, 3.10058108e-04, 0.00000000e+00,
                 0.00000000e+00, 2.27421721e-03, 1.05632062e-03, 5.42819222e-03,
                 3.48484864e-02, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
                 0.00000000e+00, 1.90858845e-05, 9.09490682e-06, 1.31265147e-04,
                 5.29163902e-06, 1.11952381e-04, 4.78452431e-06, 3.43582863e-04,
                 5.57068428e-06, 1.07167376e-03, 3.99179008e-03, 4.07753410e-03,
                 1.05749617e-04, 2.76528927e-03, 2.59244312e-05, 3.51888176e-04,
                 2.31519337e-03, 1.99211177e-03, 4.02034629e-03, 2.03778082e-04,
                 1.13483313e-02, 9.02551628e-04, 1.58182497e-03, 4.63243398e-05,
                 2.92071004e-04, 3.11923094e-05, 1.56873538e-04, 2.87205987e-05,
                 3.80543083e-05, 2.55045807e-04, 1.66878572e-04, 2.10341792e-04,
                 1.26024842e-04, 9.40663015e-05])
```

```
In [256]: 1 len (ideal_model.feature_importances_)
```

```
Out[256]: 102
```

In [296]:


```
1  # 1. helper function for plotting feature importances
2  def plot_feature_importance(ideal_model, x_train):
3      importances = ideal_model.feature_importances_
4      indices = np.argsort(importances)[: : -1]
5
6      plt.figure(figsize=(10, 8))
7      plt.bar(range(x_train.shape[1]), importances[indices])
8      plt.xticks(range(x_train.shape[1]), x_train.columns[indices], rotation = 90)
9      plt.title('Feature Importances')
10     plt.show()
11
```

In [297]:

1	<code>plot_feature_importance(ideal_model, x_train)</code>
---	--



### Feature Importances



In [316]:

```


1 # 2. Helper function for plotting feature importances (Top 20 features)
2 def plot_feature_importances(ideal_model, x_train):
3     importances = ideal_model.feature_importances_
4     indices = np.argsort(importances)[::-1][:20]
5
6     plt.figure(figsize=(10, 8))
7     plt.bar(range(len(indices)), importances[indices])
8     plt.xticks(range(len(indices)), x_train.columns[indices], rotation = 90)
9     plt.title('Feature Importances')
10    plt.xticks(rotation=90)
11    plt.show()
12

```



```
In [317]: 1 plot_feature_importances(ideal_model, x_train)
```





In [325]:

```

1 # 3. helper function for plotting feature importances
2 def plot_features(columns, importances, n=20):
3     df = (pd.DataFrame({'features': columns,
4                         'feature_importances': importances})
5           .sort_values('feature_importances', ascending=False)
6           .reset_index(drop=True))
7     # plot the dataframe
8     fig, ax = plt.subplots()
9     ax.barh(df['features'][:n], df['feature_importances'][:n])
10    ax.set_ylabel('Features')
11    ax.set_xlabel('Feature importance')
12    ax.invert_yaxis()


```

In [326]:

```

1 plot_features(x_train.columns, ideal_model.feature_importances_)

```



**Question to finish:** Why might knowing the feature importances of a trained machine learning model be helpful?

**Final challenge** what other machine learning models could you try on our dataset?

In [ ]:

1