

```
In [1]: 1 import numpy as np
        2 import matplotlib.pyplot as plt
        3 import sklearn
        4 import pandas as pd
        5 from sklearn.preprocessing import StandardScaler
        6 from sklearn.model_selection import train_test_split
        7 from sklearn import svm
        8 from sklearn.metrics import accuracy_score
```

## Data Collection and Analysis

```
In [2]: 1 data = pd.read_csv('diabetes.csv')
```

```
In [3]: 1 data
```

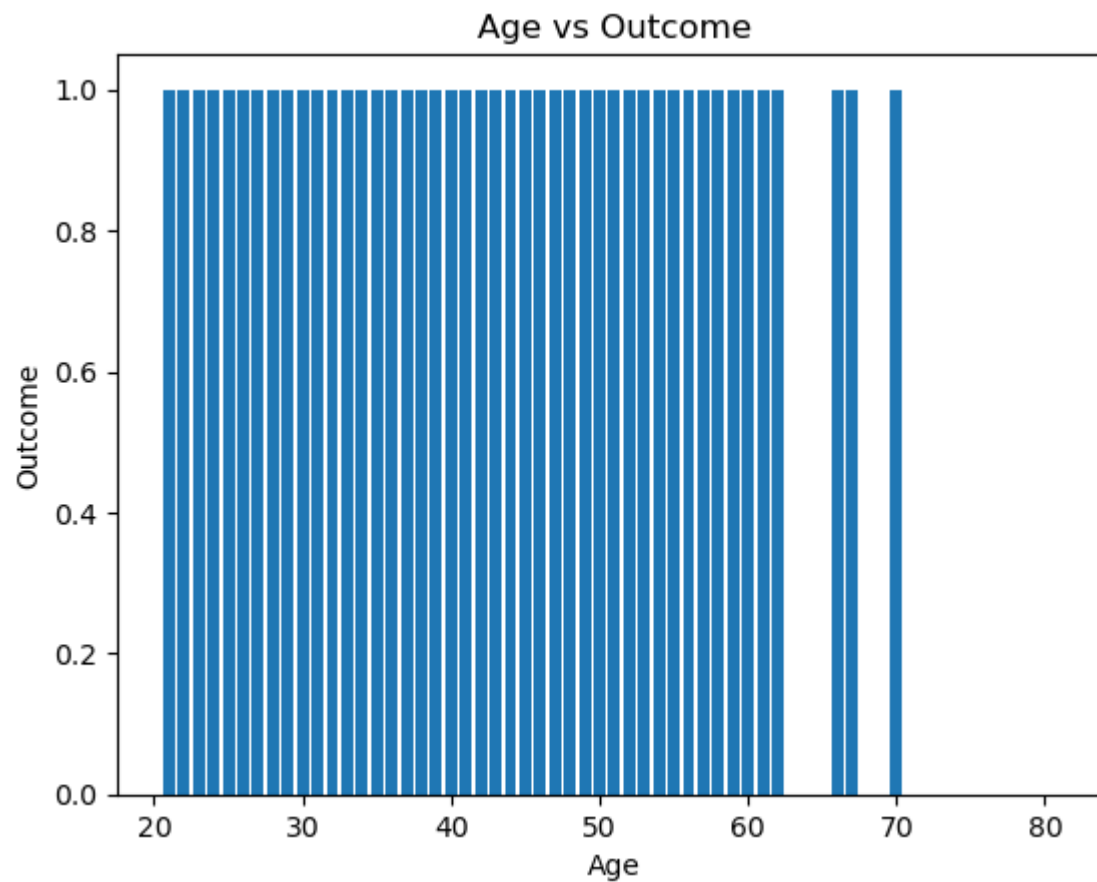
Out[3]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1
...	...	...	...	...	...	...	...	...	...
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

768 rows × 9 columns


In [4]:

```
1 # Create a bar graph of Age vs Outcome
2 plt.bar(data['Age'], data['Outcome'])
3
4 # Add Labels and title
5 plt.xlabel('Age')
6 plt.ylabel('Outcome')
7 plt.title('Age vs Outcome')
8
9 # Display the plot
10 plt.show()
```




```
In [8]: 1 import seaborn as sns
        2 # Calculate the correlation matrix
        3 corr_matrix = data.corr()
        4
        5 # Plot the correlation matrix heatmap
        6 plt.figure(figsize=(10, 8))
        7 sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
        8 plt.title('Correlation Matrix Heatmap')
        9 plt.show()
```







```
In [11]: 1 import seaborn as sns
2 from sklearn.metrics import confusion_matrix
3
4 # Extract the 'Outcome' column as the predicted labels
5 outcome = data['Outcome'].values
6
7 # Create an array of the same length with the true labels (assuming all are 0)
8 true_labels = np.zeros(len(outcome))
9
10 # Calculate the confusion matrix
11 cm = confusion_matrix(true_labels, outcome)
12
13 # Create a DataFrame for the confusion matrix
14 cm_df = pd.DataFrame(cm, index=['Actual 0', 'Actual 1'], columns=['Predicted 0', 'Predicted 1'])
15
16 # Create a heatmap using seaborn
17 plt.figure(figsize=(6, 4))
18 sns.heatmap(cm_df, annot=True, fmt='d', cmap='Blues')
19
20 # Add labels, title, and ticks
21 plt.xlabel('Predicted')
22 plt.ylabel('Actual')
23 plt.title('Confusion Matrix')
24 plt.xticks([0.5, 1.5], ['0', '1'])
25 plt.yticks([0.5, 1.5], ['0', '1'])
26 plt.show()
27
28
29 # This is still empty confusion matrix
```




```
In [12]: 1 # Use crosstab to compare 'Age' and 'Outcome' columns  
2 cross_tab = pd.crosstab(data['Age'], data['Outcome'])  
3  
4 print(cross_tab[:10])
```

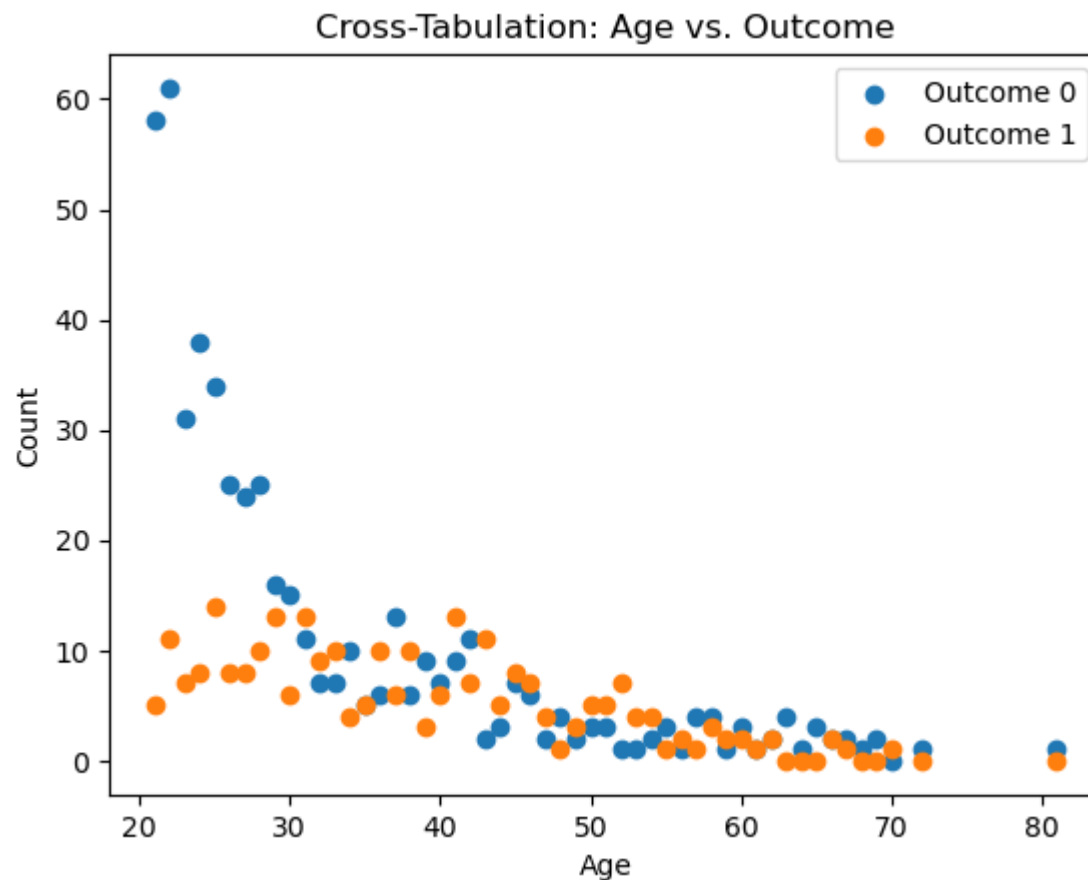
Outcome	0	1
Age		
21	58	5
22	61	11
23	31	7
24	38	8
25	34	14
26	25	8
27	24	8
28	25	10
29	16	13
30	15	6



```
In [13]: 1 # Plotting the bar chart
2 cross_tab[:20].plot(kind='bar', stacked=True)
3
4 # Adding Labels and title
5 plt.xlabel('Age')
6 plt.ylabel('Count')
7 plt.title('Comparison of Age and Outcome')
8
9 # Display the plot
10 plt.show()
```



```
In [14]: 1 # Use crosstab to compare 'Age' and 'Outcome' columns
2 cross_tab = pd.crosstab(data['Age'], data['Outcome'])
3
4 # Visualize the cross-tabulation as scatter plots
5 for outcome in cross_tab.columns:
6     plt.scatter(cross_tab.index, cross_tab[outcome], label=f'Outcome {outcome}')
7
8 # Add Labels and title
9 plt.xlabel('Age')
10 plt.ylabel('Count')
11 plt.title('Cross-Tabulation: Age vs. Outcome')
12 plt.legend()
13
14 # Display the plot
15 plt.show()
```



```
In [15]: 1 data.isna().sum()
```

```
Out[15]: Pregnancies      0
          Glucose          0
          BloodPressure    0
          SkinThickness     0
          Insulin           0
          BMI               0
          DiabetesPedigreeFunction  0
          Age              0
          Outcome           0
          dtype: int64
```

```
In [16]: 1 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies           768 non-null   int64
1   Glucose               768 non-null   int64
2   BloodPressure         768 non-null   int64
3   SkinThickness         768 non-null   int64
4   Insulin               768 non-null   int64
5   BMI                   768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                   768 non-null   int64
8   Outcome               768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [17]: 1 data.columns
```

```
Out[17]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
               'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
              dtype='object')
```

```
In [18]: 1 data.shape
```

```
Out[18]: (768, 9)
```

```
In [19]: 1 # Getting the statistical measures of the data
        2 data.describe()
```

Out[19]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240885	0.348581
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760232	0.476104
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	1.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

```
In [20]: 1 data['Outcome'].value_counts()
```

Out[20]: 0 500  
1 268  
Name: Outcome, dtype: int64

## 0 = Non-Diabetic

## 1 = Diabetic

```
In [21]: 1 data.groupby('Outcome').mean()
```

Out[21]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
Outcome								
0	3.298000	109.980000	68.184000	19.664000	68.792000	30.304200	0.429734	31.190000
1	4.865672	141.257463	70.824627	22.164179	100.335821	35.142537	0.550500	37.067164

## splitting our data

```
In [22]: 1 x = data.drop('Outcome', axis=1)
          2 y = data['Outcome']
```

```
In [23]: 1 x.shape, y.shape
```

```
Out[23]: ((768, 8), (768,))
```

```
In [24]: 1 x
```

```
Out[24]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	6	148	72	35	0	33.6	0.627	50
1	1	85	66	29	0	26.6	0.351	31
2	8	183	64	0	0	23.3	0.672	32
3	1	89	66	23	94	28.1	0.167	21
4	0	137	40	35	168	43.1	2.288	33
...	...	...	...	...	...	...	...	...
763	10	101	76	48	180	32.9	0.171	63
764	2	122	70	27	0	36.8	0.340	27
765	5	121	72	23	112	26.2	0.245	30
766	1	126	60	0	0	30.1	0.349	47
767	1	93	70	31	0	30.4	0.315	23

768 rows × 8 columns

In [25]:

1	y
---	---

Out[25]:

0	1
1	0
2	1
3	0
4	1
..	
763	0
764	0
765	0
766	1
767	0

Name: Outcome, Length: 768, dtype: int64

## Data Standardization

In [26]:

1	scalar = StandardScaler()
2	scalar.fit(x)
3	

Out[26]: StandardScaler()

In [27]:

1	standardized_data = scalar.transform(x)
---	---

```
In [28]: 1 standardized_data
```

```
Out[28]: array([[ 0.63994726,  0.84832379,  0.14964075, ...,  0.20401277,  
                0.46849198,  1.4259954 ],  
              [-0.84488505, -1.12339636, -0.16054575, ..., -0.68442195,  
                -0.36506078, -0.19067191],  
              [ 1.23388019,  1.94372388, -0.26394125, ..., -1.10325546,  
                0.60439732, -0.10558415],  
              ...,  
              [ 0.3429808 ,  0.00330087,  0.14964075, ..., -0.73518964,  
                -0.68519336, -0.27575966],  
              [-0.84488505,  0.1597866 , -0.47073225, ..., -0.24020459,  
                -0.37110101,  1.17073215],  
              [-0.84488505, -0.8730192 ,  0.04624525, ..., -0.20212881,  
                -0.47378505, -0.87137393]])
```

```
In [29]: 1 x = standardized_data  
        2 y = data['Outcome']
```

In [30]: 1 x, y

```
Out[30]: (array([[ 0.63994726,  0.84832379,  0.14964075, ...,  0.20401277,
                  0.46849198,  1.4259954 ],
                 [-0.84488505, -1.12339636, -0.16054575, ..., -0.68442195,
                  -0.36506078, -0.19067191],
                 [ 1.23388019,  1.94372388, -0.26394125, ..., -1.10325546,
                  0.60439732, -0.10558415],
                 ...,
                 [ 0.3429808 ,  0.00330087,  0.14964075, ..., -0.73518964,
                  -0.68519336, -0.27575966],
                 [-0.84488505,  0.1597866 , -0.47073225, ..., -0.24020459,
                  -0.37110101,  1.17073215],
                 [-0.84488505, -0.8730192 ,  0.04624525, ..., -0.20212881,
                  -0.47378505, -0.87137393]]),
         0      1
         1      0
         2      1
         3      0
         4      1
         ..
        763    0
        764    0
        765    0
        766    1
        767    0
        Name: Outcome, Length: 768, dtype: int64)
```

## Train Test Split

In [31]: 1 x\_train, x\_test, y\_train, y\_test = train\_test\_split(x, y, test\_size=0.2, stratify=y, random\_state=2)

In [32]: 1 x\_train.shape, x\_test.shape, y\_train.shape, y\_test.shape

Out[32]: ((614, 8), (154, 8), (614,), (154,))

## Training Model



```
In [33]: 1 model = svm.SVC(kernel='linear')
```

```
In [34]: 1 model.fit(x_train, y_train)
```

```
Out[34]: SVC(kernel='linear')
```

## Model evaluation

### 1

```
In [35]: 1 # Assuming you have obtained the predictions using model.predict(x_test)
2 from sklearn import metrics
3 y_pred = model.predict(x_test)
4
5 # Compute accuracy
6 accuracy = metrics.accuracy_score(y_test, y_pred)
7 print("Accuracy:", accuracy)
8
9 # Compute precision, recall, and F1-score
10 precision = metrics.precision_score(y_test, y_pred)
11 recall = metrics.recall_score(y_test, y_pred)
12 f1_score = metrics.f1_score(y_test, y_pred)
13
14 print("Precision:", precision)
15 print("Recall:", recall)
16 print("F1-Score:", f1_score)
```

```
Accuracy: 0.7727272727272727
Precision: 0.7567567567567568
Recall: 0.5185185185185185
F1-Score: 0.6153846153846154
```

### 2

In [36]:

```
1 from sklearn import svm
2 from sklearn.model_selection import GridSearchCV
3
4 # Define the parameter grid
5 param_grid = {
6     'C': [0.1, 1, 10],
7     'kernel': ['linear', 'rbf'],
8     'gamma': [0.1, 1, 10]
9 }
10
11 # Create the SVM model
12 model = svm.SVC()
13
14 # Create the GridSearchCV object
15 grid_search = GridSearchCV(model, param_grid, scoring='accuracy', cv=5)
16
17 # Perform grid search to find the best hyperparameters
18 grid_search.fit(x_train, y_train)
19
20 # Get the best hyperparameters
21 best_params = grid_search.best_params_
22 print("Best Hyperparameters:", best_params)
23
24 # Use the best model for prediction
25 best_model = grid_search.best_estimator_
26 y_pred = best_model.predict(x_test)
27
28 # Evaluate the best model
29 accuracy = metrics.accuracy_score(y_test, y_pred)
30 precision = metrics.precision_score(y_test, y_pred)
31 recall = metrics.recall_score(y_test, y_pred)
32 f1_score = metrics.f1_score(y_test, y_pred)
33
34 print("Accuracy:", accuracy)
35 print("Precision:", precision)
36 print("Recall:", recall)
37 print("F1-Score:", f1_score)
38
```

Best Hyperparameters: {'C': 1, 'gamma': 0.1, 'kernel': 'linear'}  
 Accuracy: 0.7727272727272727  
 Precision: 0.7567567567567568  
 Recall: 0.5185185185185185  
 F1-Score: 0.6153846153846154


In [37]: 1 y\_pred

Out[37]: array([0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0,  
 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1,  
 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,  
 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,  
 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1,  
 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0],  
 dtype=int64)

In [38]: 1 import pandas as pd  
 2  
 3 # Create a DataFrame with y\_test and y\_pred  
 4 df = pd.DataFrame({'Actual': y\_test, 'Predicted': y\_pred})  
 5  
 6 # Create the cross-tabulation  
 7 cross\_tab = pd.crosstab(df['Actual'], df['Predicted'])  
 8  
 9 print(cross\_tab)  
 10  
 11 plt.show()

Predicted	0	1
Actual		
0	91	9
1	26	28

```
In [39]: 1 # Calculate the confusion matrix
2 cm = confusion_matrix(y_test, y_pred)
3
4 # Plot the confusion matrix using a heatmap
5 plt.figure(figsize=(8, 6))
6 sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
7 plt.title('Confusion Matrix')
8 plt.xlabel('Predicted')
9 plt.ylabel('Actual')
10 plt.show()
```



### 3

## This is not a good practice:

it is not a good practice to evaluate the accuracy of your model by predicting on the training data itself. The purpose of splitting the dataset into training and testing sets is to evaluate the model's performance on unseen data.

By predicting on the training data ( `x_train` ) and comparing the predictions ( `x_train_preds` ) with the corresponding true labels ( `y_train` ), you are essentially evaluating how well your model fits the training data it has already seen. This does not provide a reliable measure of how well your model will generalize to new, unseen data.

Instead, you should use the separate testing data ( `x_test` and `y_test` ) to evaluate the performance of your model. After training your model using `model.fit(x_train, y_train)` , you can then use `model.predict(x_test)` to obtain the predicted labels for the testing data. You can then compare these predictions with the true labels ( `y_test` ) to calculate the accuracy on the testing set. This will give you a better indication of how well your model is likely to perform on new, unseen data.

## A predictive system can be built with this, based on available data;

```
In [55]: 1 model.fit(x_train, y_train)
          2 x_train_preds = model.predict(x_train)
          3 training_data_accuracy = accuracy_score(x_train_preds, y_train)
```

```
In [56]: 1 x_test_preds = model.predict(x_test)
          2 test_data_accuracy = accuracy_score(x_test_preds, y_test)
```

## Making a predictive system

```
In [69]: 1 input_data = (10,139,80,0,0,27.1,1.441,57,)
2
3 # change the input_data to numpy array
4 input_data_as_numpy_array = np.asarray(input_data)
5
6 # reshape the np array as we are predicting for one instance
7 input_data_reshape = input_data_as_numpy_array.reshape(1, -1)
8
9 std = scalar.transform(input_data_reshape)
10 print(std)
11
12
13 prediction = model.predict(std)
14 prediction
15
16 if prediction[0] == 1:
17     print('Diabetics')
18 else:
19     print('Non-Diabetics')
20     import warnings
21
22 # Ignore all warnings
23 warnings.filterwarnings("ignore")
24
```

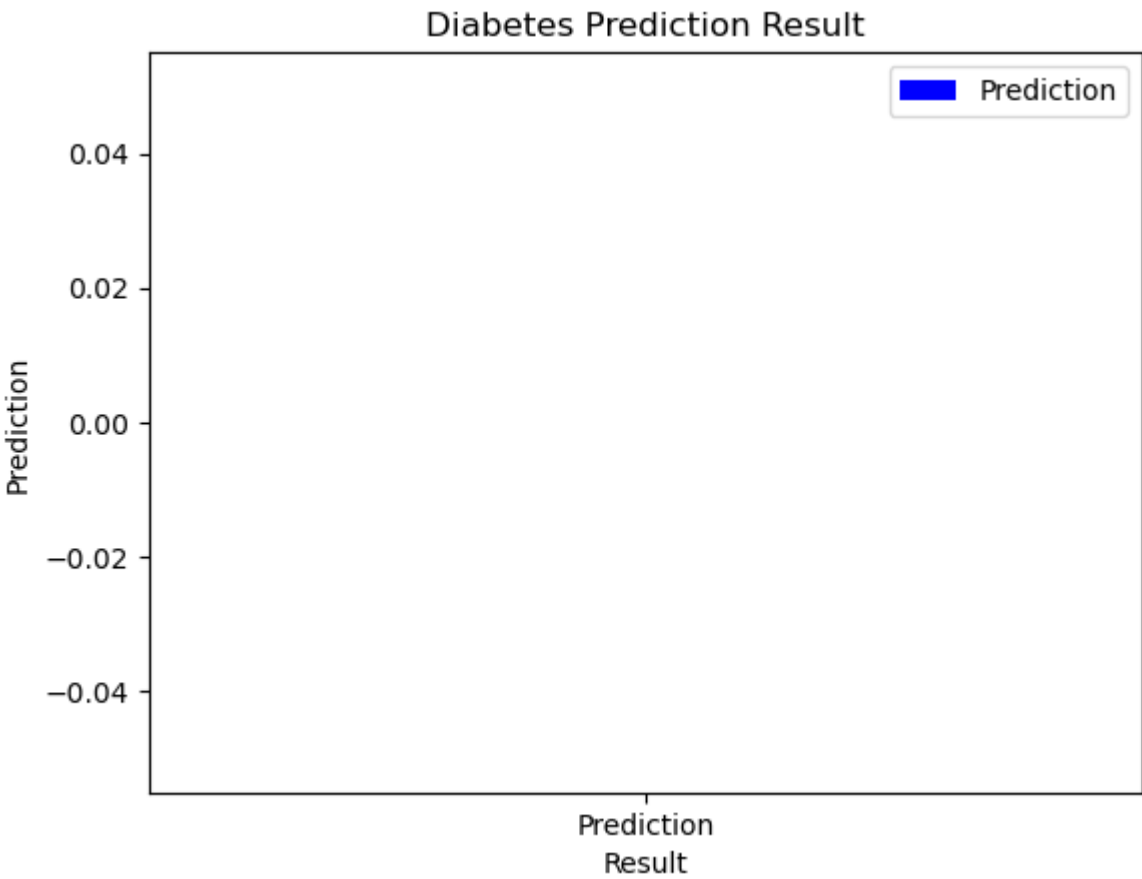
[[ 1.82781311 0.56664949 0.56322275 -1.28821221 -0.69289057 -0.62096232  
 2.92686858 2.02160968]]

Non-Diabetics

## Empty chart means 'Non-Diabetic'

```
In [80]: 1 input_data = (10, 139, 80, 0, 0, 27.1, 1.441, 57)
2
3 # Convert input_data to a numpy array
4 input_data_as_numpy_array = np.asarray(input_data)
5
6 # Reshape the np array as we are predicting for one instance
7 input_data_reshape = input_data_as_numpy_array.reshape(1, -1)
8
9 # Scale the data
10 scaler = StandardScaler()
11 std = scaler.fit_transform(input_data_reshape)
12
13 # Assuming you have the predicted result stored in the 'prediction' variable
14 prediction = 0 # Replace with your actual prediction value
15
16 # Determine the result Label
17 if prediction == 1:
18     result = 'Diabetic'
19 else:
20     result = 'Non-Diabetic'
21
22 # Plotting the result
23 plt.bar(['Prediction'], [prediction], color='blue', label='Prediction')
24 plt.xlabel('Result')
25 plt.ylabel('Prediction')
26 plt.title('Diabetes Prediction Result')
27 plt.legend()
28 plt.show()
```






In [68]:

```
1 input_data = (1,189,60,23,846,30.1,0.398,59)
2
3 # change the input_data to numpy array
4 input_data_as_numpy_array = np.asarray(input_data)
5
6 # reshape the np array as we are predicting for one instance
7 input_data_reshape = input_data_as_numpy_array.reshape(1, -1)
8
9 std = scalar.transform(input_data_reshape)
10 print(std)
11
12
13 prediction = model.predict(std)
14 prediction
15
16 if prediction[0] == 1:
17     print('Diabetics')
18 else:
19     print('Non-Diabetics')
20     import warnings
21
22 # Ignore all warnings
23 warnings.filterwarnings("ignore")
24
25
26
```

[[ -0.84488505 2.13150675 -0.47073225 0.15453319 6.65283938 -0.24020459  
 -0.2231152 2.19178518]]

Diabetics

```
In [79]: 1 input_data = (1,189,60,23,846,30.1,0.398,59)
2
3 # Convert input_data to a numpy array
4 input_data_as_numpy_array = np.asarray(input_data)
5
6 # Reshape the np array as we are predicting for one instance
7 input_data_reshape = input_data_as_numpy_array.reshape(1, -1)
8
9 # Scale the data
10 scaler = StandardScaler()
11 std = scaler.fit_transform(input_data_reshape)
12
13 # Assuming you have the predicted result stored in the 'prediction' variable
14 prediction = 1 # Replace with your actual prediction value
15
16 # Determine the result label
17 if prediction == 1:
18     result = 'Diabetic'
19 else:
20     result = 'Non-Diabetic'
21
22 # Plotting the result
23 plt.bar(['Prediction'], [prediction], color='blue', label='Prediction')
24 plt.xlabel('Result')
25 plt.ylabel('Prediction')
26 plt.title('Diabetes Prediction Result')
27 plt.legend()
28 plt.show()
```



1	
---	--