

Data description

MSSubClass: Identifies the type of dwelling involved in the sale.

20	1-STORY 1946 & NEWER ALL STYLES
30	1-STORY 1945 & OLDER
40	1-STORY W/FINISHED ATTIC ALL AGES
45	1-1/2 STORY - UNFINISHED ALL AGES
50	1-1/2 STORY FINISHED ALL AGES
60	2-STORY 1946 & NEWER
70	2-STORY 1945 & OLDER
75	2-1/2 STORY ALL AGES
80	SPLIT OR MULTI-LEVEL
85	SPLIT FOYER
90	DUPLEX - ALL STYLES AND AGES
120	1-STORY PUD (Planned Unit Development) - 1946 & NEWER
150	1-1/2 STORY PUD - ALL AGES
160	2-STORY PUD - 1946 & NEWER
180	PUD - MULTILEVEL - INCL SPLIT LEV/FOYER
190	2 FAMILY CONVERSION - ALL STYLES AND AGES

MSZoning: Identifies the general zoning classification of the sale.

A	Agriculture
C	Commercial
FV	Floating Village Residential
I	Industrial
RH	Residential High Density
RL	Residential Low Density
RP	Residential Low Density Park
RM	Residential Medium Density

LotFrontage: Linear feet of street connected to property

LotArea: Lot size in square feet

Street: Type of road access to property

Grv1	Gravel
Pave	Paved

Alley: Type of alley access to property

Grv1	Gravel
Pave	Paved
NA	No alley access

LotShape: General shape of property

Reg	Regular
IR1	Slightly irregular
IR2	Moderately Irregular
IR3	Irregular

LandContour: Flatness of the property

Lv1	Near Flat/Level
Bnk	Banked - Quick and significant rise from street grade to building
HLS	Hillside - Significant slope from side to side
Low	Depression

Utilities: Type of utilities available

AllPub	All public Utilities (E,G,W,& S)
NoSewr	Electricity, Gas, and Water (Septic Tank)
NoSeWa	Electricity and Gas Only
ELO	Electricity only

LotConfig: Lot configuration

Inside	Inside lot
Corner	Corner lot
CulDSac	Cul-de-sac
FR2	Frontage on 2 sides of property
FR3	Frontage on 3 sides of property

LandSlope: Slope of property

Gtl	Gentle slope
Mod	Moderate Slope
Sev	Severe Slope

Neighborhood: Physical locations within Ames city limits

Blmngtn	Bloomington Heights
Blueste	Bluestem
BrDale	Briardale
BrkSide	Brookside
ClearCr	Clear Creek
CollgCr	College Creek
Crawfor	Crawford
Edwards	Edwards
Gilbert	Gilbert
IDOTRR	Iowa DOT and Rail Road
MeadowV	Meadow Village
Mitchel	Mitchell
Names	North Ames
NoRidge	Northridge
NPkVill	Northpark Villa
NridgHt	Northridge Heights
NWAmes	Northwest Ames
OldTown	Old Town
SWISU	South & West of Iowa State University
Sawyer	Sawyer
SawyerW	Sawyer West
Somerst	Somerset
StoneBr	Stone Brook
Timber	Timberland
Veenker	Veenker

Condition1: Proximity to various conditions

Artery	Adjacent to arterial street
Feedr	Adjacent to feeder street
Norm	Normal
RRNn	Within 200' of North-South Railroad
RRAn	Adjacent to North-South Railroad
PosN	Near positive off-site feature--park, greenbelt, etc.
PosA	Adjacent to postive off-site feature
RRNe	Within 200' of East-West Railroad
RRAe	Adjacent to East-West Railroad

Condition2: Proximity to various conditions (if more than one is present)

Artery	Adjacent to arterial street
Feedr	Adjacent to feeder street
Norm	Normal
RRNn	Within 200' of North-South Railroad
RRAn	Adjacent to North-South Railroad
PosN	Near positive off-site feature--park, greenbelt, etc.
PosA	Adjacent to postive off-site feature
RRNe	Within 200' of East-West Railroad
RRAe	Adjacent to East-West Railroad

BldgType: Type of dwelling

1Fam	Single-family Detached
2FmCon	Two-family Conversion; originally built as one-family dwelling
Duplx	Duplex
TwnhsE	Townhouse End Unit
TwnhsI	Townhouse Inside Unit

HouseStyle: Style of dwelling

1Story	One story
1.5Fin	One and one-half story: 2nd level finished
1.5Unf	One and one-half story: 2nd level unfinished
2Story	Two story
2.5Fin	Two and one-half story: 2nd level finished
2.5Unf	Two and one-half story: 2nd level unfinished
SFoyer	Split Foyer
SLvl	Split Level

OverallQual: Rates the overall material and finish of the house

10	Very Excellent
9	Excellent
8	Very Good
7	Good
6	Above Average
5	Average
4	Below Average
3	Fair
2	Poor
1	Very Poor

OverallCond: Rates the overall condition of the house

10	Very Excellent
9	Excellent
8	Very Good
7	Good
6	Above Average
5	Average
4	Below Average
3	Fair
2	Poor
1	Very Poor

YearBuilt: Original construction date

YearRemodAdd: Remodel date (same as construction date if no remodeling or additions)

RoofStyle: Type of roof

Flat	Flat
Gable	Gable
Gambrel	Gabrel (Barn)
Hip	Hip
Mansard	Mansard
Shed	Shed

RoofMatl: Roof material

ClyTile	Clay or Tile
CompShg	Standard (Composite) Shingle
Membran	Membrane
Metal	Metal
Roll	Roll
Tar&Grv	Gravel & Tar
WdShake	Wood Shakes
WdShngl	Wood Shingles

Exterior1st: Exterior covering on house

AsbShng	Asbestos Shingles
AsphShn	Asphalt Shingles
BrkComm	Brick Common
BrkFace	Brick Face
CBlock	Cinder Block
CemntBd	Cement Board
HdBoard	Hard Board
ImStucc	Imitation Stucco
MetalSd	Metal Siding
Other	Other
Plywood	Plywood
PreCast	PreCast
Stone	Stone
Stucco	Stucco
VinylSd	Vinyl Siding
Wd Sdng	Wood Siding
WdShing	Wood Shingles

Exterior2nd: Exterior covering on house (if more than one material)

AsbShng	Asbestos Shingles
AsphShn	Asphalt Shingles
BrkComm	Brick Common
BrkFace	Brick Face
CBlock	Cinder Block
CemntBd	Cement Board
HdBoard	Hard Board
ImStucc	Imitation Stucco
MetalSd	Metal Siding
Other	Other
Plywood	Plywood
PreCast	PreCast
Stone	Stone
Stucco	Stucco
VinylSd	Vinyl Siding
Wd Sdng	Wood Siding
WdShing	Wood Shingles

MasVnrType: Masonry veneer type

BrkCmn	Brick Common
BrkFace	Brick Face
CBlock	Cinder Block
None	None
Stone	Stone

MasVnrArea: Masonry veneer area in square feet

ExterQual: Evaluates the quality of the material on the exterior

Ex	Excellent
Gd	Good
TA	Average/Typical
Fa	Fair
Po	Poor

ExterCond: Evaluates the present condition of the material on the exterior

Ex	Excellent
Gd	Good
TA	Average/Typical
Fa	Fair
Po	Poor

Foundation: Type of foundation

BrkTil	Brick & Tile
CBlock	Cinder Block
PConc	Poured Contrete
Slab	Slab
Stone	Stone
Wood	Wood

BsmtQual: Evaluates the height of the basement

Ex	Excellent (100+ inches)
Gd	Good (90-99 inches)
TA	Typical (80-89 inches)
Fa	Fair (70-79 inches)
Po	Poor (<70 inches)
NA	No Basement

BsmtCond: Evaluates the general condition of the basement

Ex	Excellent
Gd	Good
TA	Typical - slight dampness allowed
Fa	Fair - dampness or some cracking or settling
Po	Poor - Severe cracking, settling, or wetness
NA	No Basement

BsmtExposure: Refers to walkout or garden level walls

Gd	Good Exposure
Av	Average Exposure (split levels or foyers typically score average or above)
Mn	Minimum Exposure
No	No Exposure
NA	No Basement

BsmtFinType1: Rating of basement finished area

GLQ	Good Living Quarters
ALQ	Average Living Quarters
BLQ	Below Average Living Quarters
Rec	Average Rec Room
LwQ	Low Quality
Unf	Unfinished
NA	No Basement

BsmtFinSF1: Type 1 finished square feet

BsmtFinType2: Rating of basement finished area (if multiple types)

GLQ	Good Living Quarters
ALQ	Average Living Quarters
BLQ	Below Average Living Quarters
Rec	Average Rec Room
LwQ	Low Quality
Unf	Unfinished
NA	No Basement

BsmtFinSF2: Type 2 finished square feet

BsmtUnfSF: Unfinished square feet of basement area

TotalBsmtSF: Total square feet of basement area

Heating: Type of heating

Floor	Floor Furnace
GasA	Gas forced warm air furnace
GasW	Gas hot water or steam heat
Grav	Gravity furnace
OthW	Hot water or steam heat other than gas
Wall	Wall furnace

HeatingQC: Heating quality and condition

Ex	Excellent
Gd	Good
TA	Average/Typical
Fa	Fair
Po	Poor

CentralAir: Central air conditioning

N	No
Y	Yes

Electrical: Electrical system

SBrkr	Standard Circuit Breakers & Romex
FuseA	Fuse Box over 60 AMP and all Romex wiring (Average)
FuseF	60 AMP Fuse Box and mostly Romex wiring (Fair)
FuseP	60 AMP Fuse Box and mostly knob & tube wiring (poor)
Mix	Mixed

1stFlrSF: First Floor square feet

2ndFlrSF: Second floor square feet

LowQualFinSF: Low quality finished square feet (all floors)

GrLivArea: Above grade (ground) living area square feet

BsmtFullBath: Basement full bathrooms

BsmtHalfBath: Basement half bathrooms

FullBath: Full bathrooms above grade

HalfBath: Half baths above grade

Bedroom: Bedrooms above grade (does NOT include basement bedrooms)

Kitchen: Kitchens above grade

KitchenQual: Kitchen quality

Ex	Excellent
Gd	Good
TA	Typical/Average
Fa	Fair
Po	Poor

TotRmsAbvGrd: Total rooms above grade (does not include bathrooms)

Functional: Home functionality (Assume typical unless deductions are warranted)

Typ	Typical Functionality
Min1	Minor Deductions 1
Min2	Minor Deductions 2
Mod	Moderate Deductions
Maj1	Major Deductions 1
Maj2	Major Deductions 2
Sev	Severely Damaged
Sal	Salvage only

Fireplaces: Number of fireplaces

FireplaceQu: Fireplace quality

Ex	Excellent - Exceptional Masonry Fireplace
Gd	Good - Masonry Fireplace in main level
TA	Average - Prefabricated Fireplace in main living area or Masonry Fireplace in basement
Fa	Fair - Prefabricated Fireplace in basement
Po	Poor - Ben Franklin Stove
NA	No Fireplace

GarageType: Garage location

2Types	More than one type of garage
Attchd	Attached to home
Basment	Basement Garage
BuiltIn	Built-In (Garage part of house - typically has room above garage)
CarPort	Car Port
Detchd	Detached from home
NA	No Garage

GarageYrBlt: Year garage was built

GarageFinish: Interior finish of the garage

Fin	Finished
RFn	Rough Finished
Unf	Unfinished
NA	No Garage

GarageCars: Size of garage in car capacity

GarageArea: Size of garage in square feet

GarageQual: Garage quality

Ex	Excellent
Gd	Good
TA	Typical/Average
Fa	Fair
Po	Poor
NA	No Garage

GarageCond: Garage condition

Ex	Excellent
Gd	Good
TA	Typical/Average
Fa	Fair
Po	Poor
NA	No Garage

PavedDrive: Paved driveway

Y	Paved
P	Partial Pavement
N	Dirt/Gravel

WoodDeckSF: Wood deck area in square feet

OpenPorchSF: Open porch area in square feet

EnclosedPorch: Enclosed porch area in square feet

3SsnPorch: Three season porch area in square feet

ScreenPorch: Screen porch area in square feet

PoolArea: Pool area in square feet

PoolQC: Pool quality

Ex	Excellent
Gd	Good
TA	Average/Typical
Fa	Fair
NA	No Pool

Fence: Fence quality

GdPrv	Good Privacy
MnPrv	Minimum Privacy
GdWo	Good Wood
MnWw	Minimum Wood/Wire

Data Cleaning


```
In [196]: import pandas as pd
```

```
In [197]: df = pd.read_csv('house-prices-advanced-regression-techniques/train.csv')
df
```

Out[197]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	PoolArea	PoolQC	Fence
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN
...
1455	1456	60	RL	62.0	7917	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN
1456	1457	20	RL	85.0	13175	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	MnPrv
1457	1458	70	RL	66.0	9042	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	GdPrv
1458	1459	20	RL	68.0	9717	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN
1459	1460	20	RL	75.0	9937	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN

1460 rows × 81 columns



```
In [198]: df.isna().sum()
```

```
Out[198]: Id                0
MSSubClass                0
MSZoning                  0
LotFrontage             259
LotArea                  0
...
MoSold                   0
YrSold                   0
SaleType                 0
SaleCondition            0
SalePrice                0
Length: 81, dtype: int64
```

In [199]: `df.columns`

```
Out[199]: Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',
                'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',
                'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType',
                'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd',
                'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',
                'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual',
                'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',
                'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating',
                'HeatingQC', 'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF',
                'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath',
                'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual',
                'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType',
                'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual',
                'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF',
                'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'PoolQC',
                'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',
                'SaleCondition', 'SalePrice'],
                dtype='object')
```

In [200]: `df['YrSold'].info()`

```
<class 'pandas.core.series.Series'>
RangeIndex: 1460 entries, 0 to 1459
Series name: YrSold
Non-Null Count  Dtype
-----
1460 non-null   int64
dtypes: int64(1)
memory usage: 11.5 KB
```

In [201]: `df['YrSold'].value_counts()`

```
Out[201]: 2009    338
          2007    329
          2006    314
          2008    304
          2010    175
          Name: YrSold, dtype: int64
```



```
In [202]: df.YrSold[: 10]
```

```
Out[202]: 0    2008  
          1    2007  
          2    2008  
          3    2006  
          4    2008  
          5    2009  
          6    2007  
          7    2009  
          8    2008  
          9    2008  
          Name: YrSold, dtype: int64
```

```
In [203]: # To check important columns
```

```
missing_percentage = df.isnull().mean() * 100
```

```
In [204]: missing_percentage
```

```
Out[204]: Id                0.000000  
          MSSubClass        0.000000  
          MSZoning          0.000000  
          LotFrontage      17.739726  
          LotArea           0.000000  
          ...  
          MoSold            0.000000  
          YrSold            0.000000  
          SaleType          0.000000  
          SaleCondition     0.000000  
          SalePrice         0.000000  
          Length: 81, dtype: float64
```

```
In [205]: # sorted_columns
```

```
sorted_columns = missing_percentage.sort_values(ascending = False)
```

```
In [206]: sorted_columns
```

```
Out[206]: PoolQC          99.520548  
MiscFeature  96.301370  
Alley        93.767123  
Fence        80.753425  
FireplaceQu  47.260274  
...  
ExterQual    0.000000  
Exterior2nd  0.000000  
Exterior1st  0.000000  
RoofMatl     0.000000  
SalePrice    0.000000  
Length: 81, dtype: float64
```

```
In [207]: for column, percentage in sorted_columns.items():  
          print(f'{column}: {percentage: .2f}% missing')
```

PoolQC: 99.52% missing
MiscFeature: 96.30% missing
Alley: 93.77% missing
Fence: 80.75% missing
FireplaceQu: 47.26% missing
LotFrontage: 17.74% missing
GarageYrBlt: 5.55% missing
GarageCond: 5.55% missing
GarageType: 5.55% missing
GarageFinish: 5.55% missing
GarageQual: 5.55% missing
BsmtFinType2: 2.60% missing
BsmtExposure: 2.60% missing
BsmtQual: 2.53% missing
BsmtCond: 2.53% missing
BsmtFinType1: 2.53% missing
MasVnrArea: 0.55% missing
MasVnrType: 0.55% missing
Electrical: 0.07% missing
Id: 0.00% missing
Functional: 0.00% missing
Fireplaces: 0.00% missing
KitchenQual: 0.00% missing
KitchenAbvGr: 0.00% missing
BedroomAbvGr: 0.00% missing
HalfBath: 0.00% missing
FullBath: 0.00% missing
BsmtHalfBath: 0.00% missing
TotRmsAbvGrd: 0.00% missing
GarageCars: 0.00% missing
GrLivArea: 0.00% missing
GarageArea: 0.00% missing
PavedDrive: 0.00% missing
WoodDeckSF: 0.00% missing
OpenPorchSF: 0.00% missing
EnclosedPorch: 0.00% missing
3SsnPorch: 0.00% missing
ScreenPorch: 0.00% missing
PoolArea: 0.00% missing
MiscVal: 0.00% missing
MoSold: 0.00% missing
YrSold: 0.00% missing
SaleType: 0.00% missing

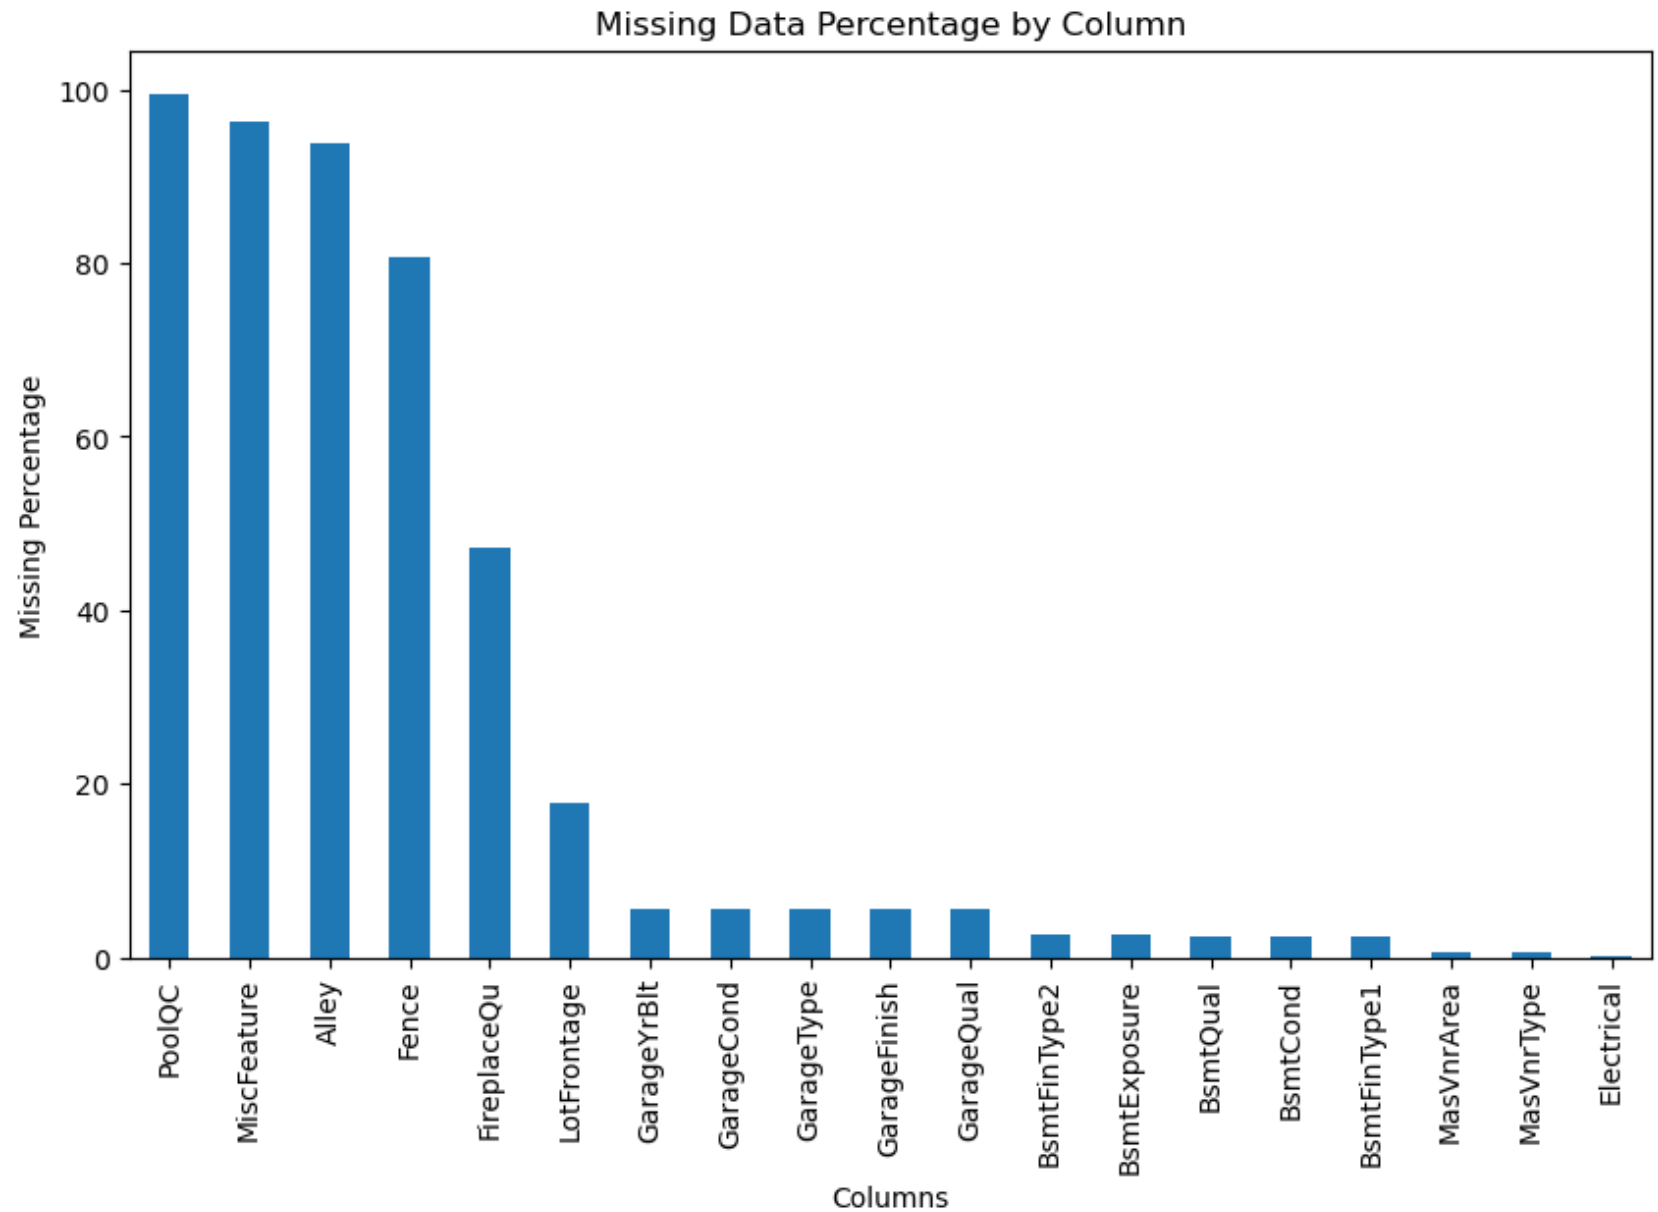
SaleCondition: 0.00% missing
BsmtFullBath: 0.00% missing
HeatingQC: 0.00% missing
LowQualFinSF: 0.00% missing
LandSlope: 0.00% missing
OverallQual: 0.00% missing
HouseStyle: 0.00% missing
BldgType: 0.00% missing
Condition2: 0.00% missing
Condition1: 0.00% missing
Neighborhood: 0.00% missing
LotConfig: 0.00% missing
YearBuilt: 0.00% missing
Utilities: 0.00% missing
LandContour: 0.00% missing
LotShape: 0.00% missing
Street: 0.00% missing
LotArea: 0.00% missing
MSZoning: 0.00% missing
OverallCond: 0.00% missing
YearRemodAdd: 0.00% missing
2ndFlrSF: 0.00% missing
BsmtFinSF2: 0.00% missing
1stFlrSF: 0.00% missing
CentralAir: 0.00% missing
MSSubClass: 0.00% missing
Heating: 0.00% missing
TotalBsmtSF: 0.00% missing
BsmtUnfSF: 0.00% missing
BsmtFinSF1: 0.00% missing
RoofStyle: 0.00% missing
Foundation: 0.00% missing
ExterCond: 0.00% missing
ExterQual: 0.00% missing
Exterior2nd: 0.00% missing
Exterior1st: 0.00% missing
RoofMatl: 0.00% missing
SalePrice: 0.00% missing

```
In [208]: # Filter columns with missing values
missing_columns = sorted_columns[sorted_columns > 0]
missing_columns
```

```
Out[208]: PoolQC          99.520548
MiscFeature       96.301370
Alley            93.767123
Fence            80.753425
FireplaceQu      47.260274
LotFrontage      17.739726
GarageYrBlt       5.547945
GarageCond        5.547945
GarageType        5.547945
GarageFinish      5.547945
GarageQual        5.547945
BsmtFinType2      2.602740
BsmtExposure      2.602740
BsmtQual          2.534247
BsmtCond          2.534247
BsmtFinType1      2.534247
MasVnrArea        0.547945
MasVnrType        0.547945
Electrical        0.068493
dtype: float64
```

```
In [209]: import matplotlib.pyplot as plt
```

```
In [210]: # visualizing missing percentage
plt.figure(figsize=(10, 6))
missing_columns.plot(kind='bar')
plt.title('Missing Data Percentage by Column')
plt.xlabel('Columns')
plt.ylabel('Missing Percentage')
plt.show()
```



These are the columns with missing values, now we know the columns to drop.

I am going to drop this columns with the highest missing values:

- PoolQC: 99.52% missing

- MiscFeature: 96.30% missing
- Alley: 93.77% missing
- Fence: 80.75% missing
- FireplaceQu: 47.26% missing

```
In [211]: # Dropping multiple columns
columns_to_drop = ['PoolQC', 'MiscFeature', 'Alley', 'Fence', 'FireplaceQu']
df = df.drop(columns_to_drop, axis=1)
```

```
In [212]: df
```

Out[212]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	LotShape	LandContour	Utilities	LotConfig	...	EnclosedPorch	3Ssn
0	1	60	RL	65.0	8450	Pave	Reg	Lvl	AllPub	Inside	...	0	
1	2	20	RL	80.0	9600	Pave	Reg	Lvl	AllPub	FR2	...	0	
2	3	60	RL	68.0	11250	Pave	IR1	Lvl	AllPub	Inside	...	0	
3	4	70	RL	60.0	9550	Pave	IR1	Lvl	AllPub	Corner	...	272	
4	5	60	RL	84.0	14260	Pave	IR1	Lvl	AllPub	FR2	...	0	
...	
1455	1456	60	RL	62.0	7917	Pave	Reg	Lvl	AllPub	Inside	...	0	
1456	1457	20	RL	85.0	13175	Pave	Reg	Lvl	AllPub	Inside	...	0	
1457	1458	70	RL	66.0	9042	Pave	Reg	Lvl	AllPub	Inside	...	0	
1458	1459	20	RL	68.0	9717	Pave	Reg	Lvl	AllPub	Inside	...	112	
1459	1460	20	RL	75.0	9937	Pave	Reg	Lvl	AllPub	Inside	...	0	

1460 rows × 76 columns



```
In [213]: df.shape
```

Out[213]: (1460, 76)

Copy our data sets

In [214]: `df = df.copy()`

In [215]: `df`

Out[215]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	LotShape	LandContour	Utilities	LotConfig	...	EnclosedPorch	3Ssn
0	1	60	RL	65.0	8450	Pave	Reg	Lvl	AllPub	Inside	...	0	
1	2	20	RL	80.0	9600	Pave	Reg	Lvl	AllPub	FR2	...	0	
2	3	60	RL	68.0	11250	Pave	IR1	Lvl	AllPub	Inside	...	0	
3	4	70	RL	60.0	9550	Pave	IR1	Lvl	AllPub	Corner	...	272	
4	5	60	RL	84.0	14260	Pave	IR1	Lvl	AllPub	FR2	...	0	
...	
1455	1456	60	RL	62.0	7917	Pave	Reg	Lvl	AllPub	Inside	...	0	
1456	1457	20	RL	85.0	13175	Pave	Reg	Lvl	AllPub	Inside	...	0	
1457	1458	70	RL	66.0	9042	Pave	Reg	Lvl	AllPub	Inside	...	0	
1458	1459	20	RL	68.0	9717	Pave	Reg	Lvl	AllPub	Inside	...	112	
1459	1460	20	RL	75.0	9937	Pave	Reg	Lvl	AllPub	Inside	...	0	

1460 rows × 76 columns



```
In [216]: df.isnull().sum()
```

```
Out[216]: Id                0
MSSubClass                0
MSZoning                  0
LotFrontage              259
LotArea                  0
...
MoSold                    0
YrSold                    0
SaleType                  0
SaleCondition             0
SalePrice                 0
Length: 76, dtype: int64
```

Filling NaN values in column 'LotFrontage' with the median value

```
In [217]: median_value = df['LotFrontage'].median()
df['LotFrontage'].fillna(median_value, inplace=True)
```

```
In [218]: df.isna().sum()
```

```
Out[218]: Id                0
MSSubClass                0
MSZoning                  0
LotFrontage              0
LotArea                  0
..
MoSold                    0
YrSold                    0
SaleType                  0
SaleCondition             0
SalePrice                 0
Length: 76, dtype: int64
```

In [219]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 76 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    1460 non-null   int64
1   MSSubClass            1460 non-null   int64
2   MSZoning              1460 non-null   object
3   LotFrontage          1460 non-null   float64
4   LotArea              1460 non-null   int64
5   Street               1460 non-null   object
6   LotShape             1460 non-null   object
7   LandContour          1460 non-null   object
8   Utilities            1460 non-null   object
9   LotConfig            1460 non-null   object
10  LandSlope            1460 non-null   object
11  Neighborhood         1460 non-null   object
12  Condition1           1460 non-null   object
13  Condition2           1460 non-null   object
14  BldgType             1460 non-null   object
15  HouseStyle           1460 non-null   object
16  OverallQual          1460 non-null   int64
17  OverallCond          1460 non-null   int64
18  YearBuilt            1460 non-null   int64
19  YearRemodAdd         1460 non-null   int64
20  RoofStyle            1460 non-null   object
21  RoofMatl            1460 non-null   object
22  Exterior1st          1460 non-null   object
23  Exterior2nd          1460 non-null   object
24  MasVnrType           1452 non-null   object
25  MasVnrArea           1452 non-null   float64
26  ExterQual            1460 non-null   object
27  ExterCond            1460 non-null   object
28  Foundation           1460 non-null   object
29  BsmtQual             1423 non-null   object
30  BsmtCond            1423 non-null   object
31  BsmtExposure         1422 non-null   object
32  BsmtFinType1         1423 non-null   object
33  BsmtFinSF1           1460 non-null   int64
34  BsmtFinType2         1422 non-null   object
35  BsmtFinSF2           1460 non-null   int64
36  BsmtUnfSF            1460 non-null   int64
37  TotalBsmtSF          1460 non-null   int64

```


```
38 Heating          1460 non-null object
39 HeatingQC         1460 non-null object
40 CentralAir        1460 non-null object
41 Electrical         1459 non-null object
42 1stFlrSF           1460 non-null int64
43 2ndFlrSF           1460 non-null int64
44 LowQualFinSF       1460 non-null int64
45 GrLivArea          1460 non-null int64
46 BsmntFullBath      1460 non-null int64
47 BsmntHalfBath      1460 non-null int64
48 FullBath           1460 non-null int64
49 HalfBath           1460 non-null int64
50 BedroomAbvGr       1460 non-null int64
51 KitchenAbvGr       1460 non-null int64
52 KitchenQual        1460 non-null object
53 TotRmsAbvGrd       1460 non-null int64
54 Functional         1460 non-null object
55 Fireplaces         1460 non-null int64
56 GarageType         1379 non-null object
57 GarageYrBlt        1379 non-null float64
58 GarageFinish       1379 non-null object
59 GarageCars         1460 non-null int64
60 GarageArea         1460 non-null int64
61 GarageQual         1379 non-null object
62 GarageCond         1379 non-null object
63 PavedDrive         1460 non-null object
64 WoodDeckSF         1460 non-null int64
65 OpenPorchSF        1460 non-null int64
66 EnclosedPorch      1460 non-null int64
67 3SsnPorch          1460 non-null int64
68 ScreenPorch        1460 non-null int64
69 PoolArea           1460 non-null int64
70 MiscVal            1460 non-null int64
71 MoSold             1460 non-null int64
72 YrSold             1460 non-null int64
73 SaleType           1460 non-null object
74 SaleCondition       1460 non-null object
75 SalePrice          1460 non-null int64
dtypes: float64(3), int64(35), object(38)
memory usage: 867.0+ KB
```

In [220]: `df.describe()`

Out[220]:

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	1452.000000
mean	730.500000	56.897260	69.863699	10516.828082	6.099315	5.575342	1971.267808	1984.865753	103.685262
std	421.610009	42.300571	22.027677	9981.264932	1.382997	1.112799	30.202904	20.645407	181.066207
min	1.000000	20.000000	21.000000	1300.000000	1.000000	1.000000	1872.000000	1950.000000	0.000000
25%	365.750000	20.000000	60.000000	7553.500000	5.000000	5.000000	1954.000000	1967.000000	0.000000
50%	730.500000	50.000000	69.000000	9478.500000	6.000000	5.000000	1973.000000	1994.000000	0.000000
75%	1095.250000	70.000000	79.000000	11601.500000	7.000000	6.000000	2000.000000	2004.000000	166.000000
max	1460.000000	190.000000	313.000000	215245.000000	10.000000	9.000000	2010.000000	2010.000000	1600.000000

8 rows × 38 columns



In [221]: `df.columns`

Out[221]: Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',
 'LotShape', 'LandContour', 'Utilities', 'LotConfig', 'LandSlope',
 'Neighborhood', 'Condition1', 'Condition2', 'BldgType', 'HouseStyle',
 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd', 'RoofStyle',
 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType', 'MasVnrArea',
 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual', 'BsmtCond',
 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1', 'BsmtFinType2',
 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating', 'HeatingQC',
 'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF',
 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath',
 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual', 'TotRmsAbvGrd',
 'Functional', 'Fireplaces', 'GarageType', 'GarageYrBlt', 'GarageFinish',
 'GarageCars', 'GarageArea', 'GarageQual', 'GarageCond', 'PavedDrive',
 'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch',
 'ScreenPorch', 'PoolArea', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',
 'SaleCondition', 'SalePrice'],
 dtype='object')

Exploratory Data Analysis

```
In [222]: # Calculate the correlation matrix
corr_matrix = df.corr()
corr_matrix.head()
```

Out[222]:

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	BsmtFin
Id	1.000000	0.011156	-0.009921	-0.033226	-0.028365	0.012609	-0.012713	-0.021998	-0.050298	-0.005
MSSubClass	0.011156	1.000000	-0.356718	-0.139781	0.032628	-0.059316	0.027850	0.040581	0.022936	-0.069
LotFrontage	-0.009921	-0.356718	1.000000	0.304522	0.234812	-0.053281	0.116685	0.083348	0.179459	0.214
LotArea	-0.033226	-0.139781	0.304522	1.000000	0.105806	-0.005636	0.014228	0.013788	0.104160	0.214
OverallQual	-0.028365	0.032628	0.234812	0.105806	1.000000	-0.091932	0.572323	0.550684	0.411876	0.239

5 rows × 38 columns



```
In [223]: import seaborn as sns
```

```
In [224]: # Select important columns for analysis
selected_columns = ['SalePrice', 'OverallQual', 'GrLivArea', 'TotalBsmtSF', 'GarageArea', 'YearBuilt']
```

```
In [225]: # Subset the dataframe with selected columns
selected_df = df[selected_columns]
```


In [226]: `selected_df.head()`

Out[226]:

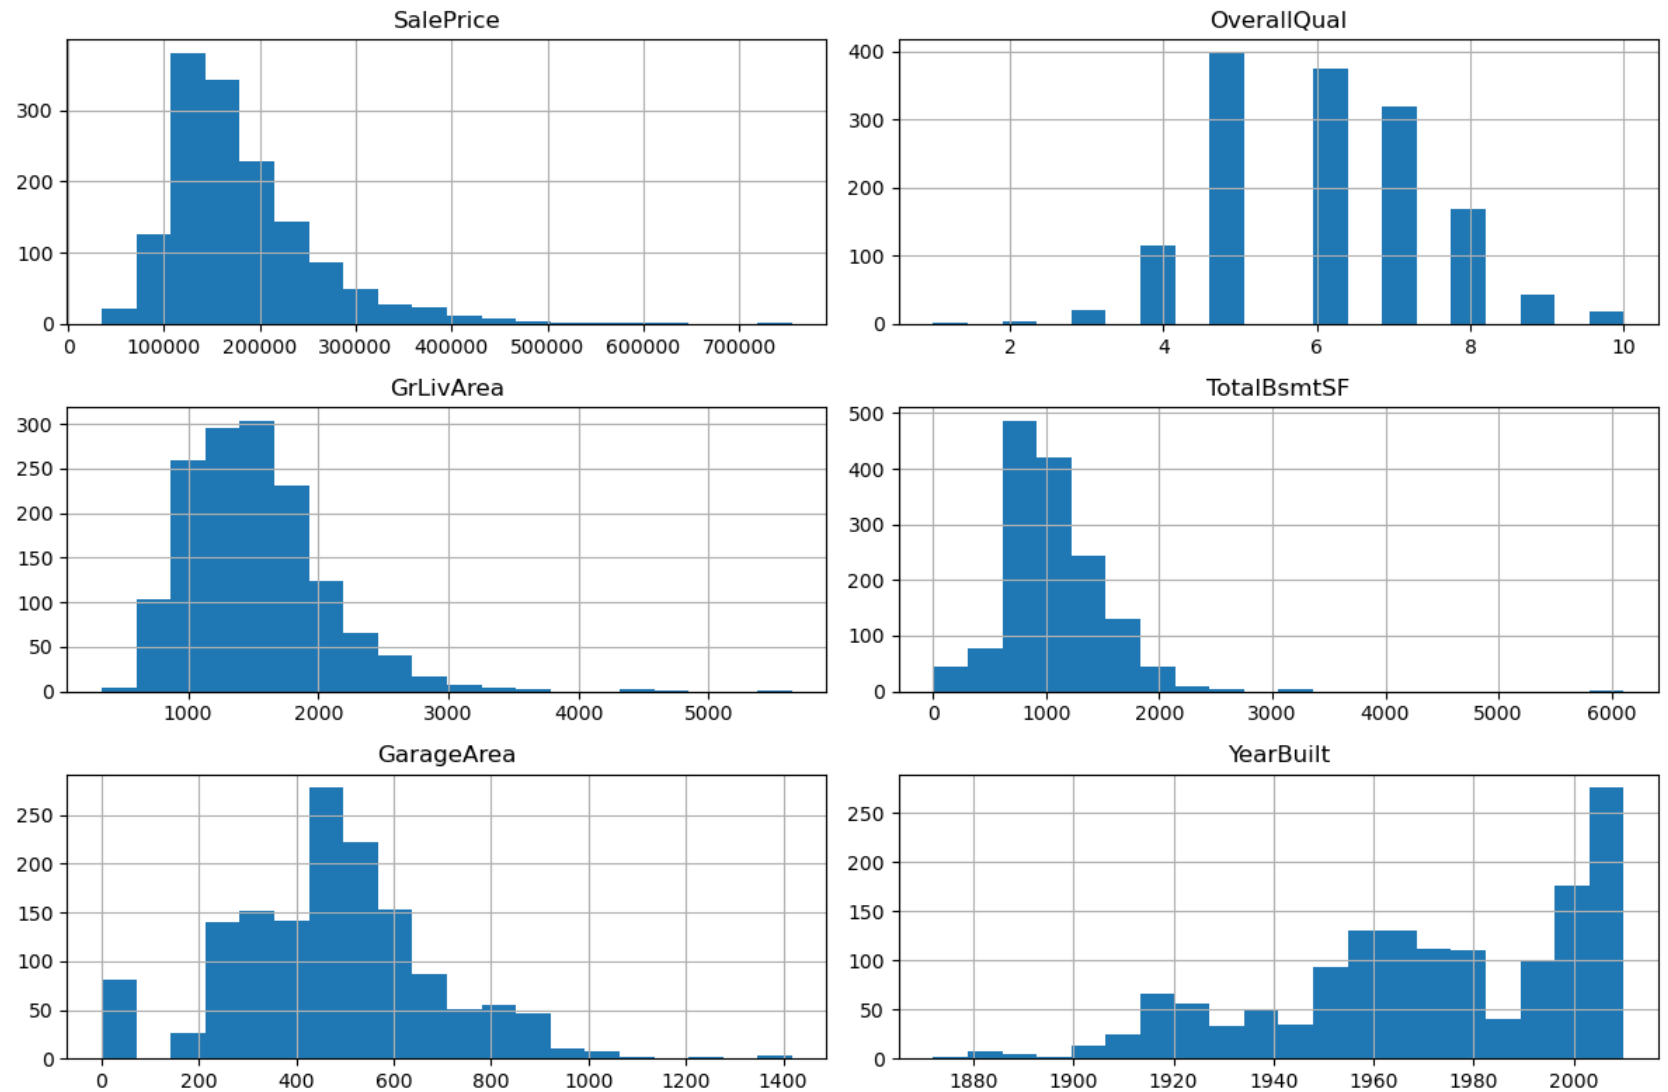
	SalePrice	OverallQual	GrLivArea	TotalBsmtSF	GarageArea	YearBuilt
0	208500	7	1710	856	548	2003
1	181500	6	1262	1262	460	1976
2	223500	7	1786	920	608	2001
3	140000	7	1717	756	642	1915
4	250000	8	2198	1145	836	2000

In [227]: `# Basic statistics of the selected columns`
`selected_statistics = selected_df.describe()`
`print(selected_statistics)`

	SalePrice	OverallQual	GrLivArea	TotalBsmtSF	GarageArea	\
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	
mean	180921.195890	6.099315	1515.463699	1057.429452	472.980137	
std	79442.502883	1.382997	525.480383	438.705324	213.804841	
min	34900.000000	1.000000	334.000000	0.000000	0.000000	
25%	129975.000000	5.000000	1129.500000	795.750000	334.500000	
50%	163000.000000	6.000000	1464.000000	991.500000	480.000000	
75%	214000.000000	7.000000	1776.750000	1298.250000	576.000000	
max	755000.000000	10.000000	5642.000000	6110.000000	1418.000000	


	YearBuilt
count	1460.000000
mean	1971.267808
std	30.202904
min	1872.000000
25%	1954.000000
50%	1973.000000
75%	2000.000000
max	2010.000000


```
In [228]: # Histograms for selected columns  
selected_df.hist(bins=20, figsize=(12, 8))  
plt.tight_layout()  
plt.show()
```



```
In [229]: # Scatter plot: SalePrice vs. GrLivArea
plt.scatter(selected_df['GrLivArea'], selected_df['SalePrice'])
plt.xlabel('GrLivArea')
plt.ylabel('SalePrice')
plt.title('SalePrice vs. GrLivArea')
plt.show()

# Scatter plot: SalePrice vs. YearBuilt
plt.scatter(selected_df['YearBuilt'], selected_df['SalePrice'])
plt.xlabel('YearBuilt')
plt.ylabel('SalePrice')
plt.title('SalePrice vs. YearBuilt')
plt.show()
```






```
In [230]: # Correlation matrix heatmap
corr_matrix = selected_df.corr()
corr_matrix
```

Out[230]:

	SalePrice	OverallQual	GrLivArea	TotalBsmntSF	GarageArea	YearBuilt
SalePrice	1.000000	0.790982	0.708624	0.613581	0.623431	0.522897
OverallQual	0.790982	1.000000	0.593007	0.537808	0.562022	0.572323
GrLivArea	0.708624	0.593007	1.000000	0.454868	0.468997	0.199010
TotalBsmntSF	0.613581	0.537808	0.454868	1.000000	0.486665	0.391452
GarageArea	0.623431	0.562022	0.468997	0.486665	1.000000	0.478954
YearBuilt	0.522897	0.572323	0.199010	0.391452	0.478954	1.000000

```
In [231]: plt.figure(figsize=(10, 8))  
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')  
plt.title('Correlation Matrix')  
plt.show()
```



Feature Engineering

```
In [232]: df_main = pd.read_csv('house-prices-advanced-regression-techniques/train.csv')
```

```
In [233]: df_main
```

Out[233]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	PoolArea	PoolQC	Fence
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN
...
1455	1456	60	RL	62.0	7917	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN
1456	1457	20	RL	85.0	13175	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	MnPrv
1457	1458	70	RL	66.0	9042	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	GdPrv
1458	1459	20	RL	68.0	9717	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN
1459	1460	20	RL	75.0	9937	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN

1460 rows × 81 columns



```
In [234]: # Dropping multiple columns
columns_to_drop = ['PoolQC', 'MiscFeature', 'Alley', 'Fence', 'FireplaceQu']
df_main = df_main.drop(columns_to_drop, axis=1)
```


In [235]: df_main

Out[235]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	LotShape	LandContour	Utilities	LotConfig	...	EnclosedPorch	3Ssn
0	1	60	RL	65.0	8450	Pave	Reg	Lvl	AllPub	Inside	...	0	
1	2	20	RL	80.0	9600	Pave	Reg	Lvl	AllPub	FR2	...	0	
2	3	60	RL	68.0	11250	Pave	IR1	Lvl	AllPub	Inside	...	0	
3	4	70	RL	60.0	9550	Pave	IR1	Lvl	AllPub	Corner	...	272	
4	5	60	RL	84.0	14260	Pave	IR1	Lvl	AllPub	FR2	...	0	
...	
1455	1456	60	RL	62.0	7917	Pave	Reg	Lvl	AllPub	Inside	...	0	
1456	1457	20	RL	85.0	13175	Pave	Reg	Lvl	AllPub	Inside	...	0	
1457	1458	70	RL	66.0	9042	Pave	Reg	Lvl	AllPub	Inside	...	0	
1458	1459	20	RL	68.0	9717	Pave	Reg	Lvl	AllPub	Inside	...	112	
1459	1460	20	RL	75.0	9937	Pave	Reg	Lvl	AllPub	Inside	...	0	

1460 rows × 76 columns



In [236]: *# due to large columns using this code may not cover the null values to see*
df_main.isnull().sum()

Out[236]:

```

Id                0
MSSubClass        0
MSZoning          0
LotFrontage      259
LotArea           0
...
MoSold            0
YrSold            0
SaleType          0
SaleCondition     0
SalePrice         0
Length: 76, dtype: int64

```

Checking for numerical columns

```
In [237]: for label, content in df_main.items():  
          if pd.api.types.is_numeric_dtype(content):  
              print(label)
```

Id
MSSubClass
LotFrontage
LotArea
OverallQual
OverallCond
YearBuilt
YearRemodAdd
MasVnrArea
BsmtFinSF1
BsmtFinSF2
BsmtUnfSF
TotalBsmtSF
1stFlrSF
2ndFlrSF
LowQualFinSF
GrLivArea
BsmtFullBath
BsmtHalfBath
FullBath
HalfBath
BedroomAbvGr
KitchenAbvGr
TotRmsAbvGrd
Fireplaces
GarageYrBlt
GarageCars
GarageArea
WoodDeckSF
OpenPorchSF
EnclosedPorch
3SsnPorch
ScreenPorch
PoolArea
MiscVal
MoSold
YrSold
SalePrice

```
In [238]: # check for which numeric columns have null values
for label, content in df_main.items():
    if pd.api.types.is_numeric_dtype(content):
        if pd.isnull(content).sum():
            print(label)
```

LotFrontage
MasVnrArea
GarageYrBlt

```
In [239]: # Fill numeric rows with the median
for label, content in df_main.items():
    if pd.api.types.is_numeric_dtype(content):
        if pd.isnull(content).sum():
            # Add a binary column which tells us if the data was missing or not
            df_main[label+'_is_missing'] = pd.isnull(content)
            # Fill missing numeric values with median
            df_main[label] = content.fillna(content.median())
```

In [240]: df_main

Out[240]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	LotShape	LandContour	Utilities	LotConfig	...	PoolArea	MiscVal	N
0	1	60	RL	65.0	8450	Pave	Reg	Lvl	AllPub	Inside	...	0	0	
1	2	20	RL	80.0	9600	Pave	Reg	Lvl	AllPub	FR2	...	0	0	
2	3	60	RL	68.0	11250	Pave	IR1	Lvl	AllPub	Inside	...	0	0	
3	4	70	RL	60.0	9550	Pave	IR1	Lvl	AllPub	Corner	...	0	0	
4	5	60	RL	84.0	14260	Pave	IR1	Lvl	AllPub	FR2	...	0	0	
...	
1455	1456	60	RL	62.0	7917	Pave	Reg	Lvl	AllPub	Inside	...	0	0	
1456	1457	20	RL	85.0	13175	Pave	Reg	Lvl	AllPub	Inside	...	0	0	
1457	1458	70	RL	66.0	9042	Pave	Reg	Lvl	AllPub	Inside	...	0	2500	
1458	1459	20	RL	68.0	9717	Pave	Reg	Lvl	AllPub	Inside	...	0	0	
1459	1460	20	RL	75.0	9937	Pave	Reg	Lvl	AllPub	Inside	...	0	0	

1460 rows × 79 columns



In [241]: df_main.isnull().sum()

Out[241]:

Id	0
MSSubClass	0
MSZoning	0
LotFrontage	0
LotArea	0
...	..
SaleCondition	0
SalePrice	0
LotFrontage_is_missing	0
MasVnrArea_is_missing	0
GarageYrBlt_is_missing	0
Length: 79, dtype: int64	

Filling and turning categorical variables into numbers

```
In [242]: # check for columns which aren't numeric  
for label, content in df_main.items():  
    if not pd.api.types.is_numeric_dtype(content):  
        print(label)
```


MSZoning
Street
LotShape
LandContour
Utilities
LotConfig
LandSlope
Neighborhood
Condition1
Condition2
BldgType
HouseStyle
RoofStyle
RoofMatl
Exterior1st
Exterior2nd
MasVnrType
ExterQual
ExterCond
Foundation
BsmtQual
BsmtCond
BsmtExposure
BsmtFinType1
BsmtFinType2
Heating
HeatingQC
CentralAir
Electrical
KitchenQual
Functional
GarageType
GarageFinish
GarageQual
GarageCond
PavedDrive
SaleType
SaleCondition

Turn categorical variables into numbers and fill missing

```
In [243]: # Turn categorical variables into numbers and fill missing
for label, content in df_main.items():
    if not pd.api.types.is_numeric_dtype(content):
        # Add binary column to indicate whether sample has missing value
        df_main[label+'_is_missing'] = pd.isnull(content)
        # Turn categories into numbers and add +1
        df_main[label] = pd.Categorical(content).codes + 1
```

```
In [244]: df_main.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Columns: 117 entries, Id to SaleCondition_is_missing
dtypes: bool(41), float64(3), int64(35), int8(38)
memory usage: 546.2 KB
```

```
In [245]: df_main.T
```

Out[245]:

	0	1	2	3	4	5	6	7	8	9	...	1450	1451	1452	1453	1454
Id	1	2	3	4	5	6	7	8	9	10	...	1451	1452	1453	1454	1455
MSSubClass	60	20	60	70	60	50	20	60	50	190	...	90	20	180	20	20
MSZoning	4	4	4	4	4	4	4	4	5	4	...	4	4	5	4	2
LotFrontage	65.0	80.0	68.0	60.0	84.0	85.0	75.0	69.0	51.0	50.0	...	60.0	78.0	35.0	90.0	62.0
LotArea	8450	9600	11250	9550	14260	14115	10084	10382	6120	7420	...	9000	9262	3675	17217	7500
...
GarageQual_is_missing	False	False	False	False	False	False	False	False	False	False	...	True	False	False	True	False
GarageCond_is_missing	False	False	False	False	False	False	False	False	False	False	...	True	False	False	True	False
PavedDrive_is_missing	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False
SaleType_is_missing	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False
SaleCondition_is_missing	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False

117 rows × 1460 columns

```
In [246]: df_main.columns
```

```
Out[246]: Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',  
               'LotShape', 'LandContour', 'Utilities', 'LotConfig',  
               ...  
               'Electrical_is_missing', 'KitchenQual_is_missing',  
               'Functional_is_missing', 'GarageType_is_missing',  
               'GarageFinish_is_missing', 'GarageQual_is_missing',  
               'GarageCond_is_missing', 'PavedDrive_is_missing', 'SaleType_is_missing',  
               'SaleCondition_is_missing'],  
              dtype='object', length=117)
```

```
In [247]: df_main.SalePrice
```

```
Out[247]: 0      208500  
          1      181500  
          2      223500  
          3      140000  
          4      250000  
          ...  
          1455     175000  
          1456     210000  
          1457     266500  
          1458     142125  
          1459     147500  
          Name: SalePrice, Length: 1460, dtype: int64
```

```
In [248]: df_main.YrSold
```

```
Out[248]: 0      2008  
          1      2007  
          2      2008  
          3      2006  
          4      2008  
          ...  
          1455     2007  
          1456     2010  
          1457     2010  
          1458     2010  
          1459     2008  
          Name: YrSold, Length: 1460, dtype: int64
```

```
In [249]: df_main.YrSold.value_counts()
```

```
Out[249]: 2009    338
          2007    329
          2006    314
          2008    304
          2010    175
          Name: YrSold, dtype: int64
```

```
In [250]: df_main.shape
```

```
Out[250]: (1460, 117)
```

Separating the data

```
In [251]: # Filter data for df_train
df_train = df_main[(df_main['YrSold'] >= 2006) & (df_main['YrSold'] <= 2009)]

# Filter data for df_val
df_val = df_main[df_main['YrSold'] == 2010]

# Output the lengths of df_val and df_train
len(df_train), len(df_val)
```

```
Out[251]: (1285, 175)
```

Splitting the dataset

```
In [252]: # split data into x and y
x_train, y_train = df_train.drop('SalePrice', axis = 1), df_train.SalePrice
x_valid, y_valid = df_val.drop('SalePrice', axis = 1), df_val.SalePrice
x_train.shape, y_train.shape, x_valid.shape, y_valid.shape
```


```
Out[252]: ((1285, 116), (1285,), (175, 116), (175,))
```

In [254]: x_train

Out[254]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	LotShape	LandContour	Utilities	LotConfig	...	Electrical_is_missing
0	1	60	4	65.0	8450	2	4	4	1	5	...	False
1	2	20	4	80.0	9600	2	4	4	1	3	...	False
2	3	60	4	68.0	11250	2	1	4	1	5	...	False
3	4	70	4	60.0	9550	2	1	4	1	1	...	False
4	5	60	4	84.0	14260	2	1	4	1	3	...	False
...
1452	1453	180	5	35.0	3675	2	4	4	1	5	...	False
1453	1454	20	4	90.0	17217	2	4	4	1	5	...	False
1454	1455	20	2	62.0	7500	2	4	4	1	5	...	False
1455	1456	60	4	62.0	7917	2	4	4	1	5	...	False
1459	1460	20	4	75.0	9937	2	4	4	1	5	...	False

1285 rows × 116 columns



In [255]: y_train

Out[255]:

```

0      208500
1      181500
2      223500
3      140000
4      250000
...
1452    145000
1453     84500
1454    185000
1455    175000
1459    147500

```

Name: SalePrice, Length: 1285, dtype: int64

In [256]: y_valid

```
Out[256]: 16      149000
          24      154000
          26      134800
          27      306000
          33      165500
          ...
          1438    149700
          1446    157900
          1456    210000
          1457    266500
          1458    142125
```


Name: SalePrice, Length: 175, dtype: int64

In [257]: x_valid

Out[257]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	LotShape	LandContour	Utilities	LotConfig	...	Electrical_is_missing
16	17	20	4	69.0	11241	2	1	4	1	2	...	False
24	25	20	4	69.0	8246	2	1	4	1	5	...	False
26	27	20	4	60.0	7200	2	4	4	1	1	...	False
27	28	20	4	98.0	11478	2	4	4	1	5	...	False
33	34	20	4	70.0	10552	2	1	4	1	5	...	False
...
1438	1439	20	5	90.0	7407	2	4	4	1	5	...	False
1446	1447	20	4	69.0	26142	2	1	4	1	2	...	False
1456	1457	20	4	85.0	13175	2	4	4	1	5	...	False
1457	1458	70	4	66.0	9042	2	4	4	1	5	...	False
1458	1459	20	4	68.0	9717	2	4	4	1	5	...	False

175 rows × 116 columns



Building an evaluation function

```
In [258]: from sklearn.metrics import mean_squared_log_error, mean_absolute_error, r2_score

def rmsle(y_test, y_preds):
    """
    Calculate root mean squared log error between predictions and true labels
    """
    return np.sqrt(mean_squared_log_error(y_test, y_preds))

def show_scores(model, x_train, y_train, x_valid, y_valid):
    train_preds = model.predict(x_train)
    val_preds = model.predict(x_valid)

    scores = {
        'Training MAE': mean_absolute_error(y_train, train_preds),
        'Valid MAE': mean_absolute_error(y_valid, val_preds),
        'Training RMSLE': rmsle(y_train, train_preds),
        'Valid RMSLE': rmsle(y_valid, val_preds),
        'Training R^2': r2_score(y_train, train_preds),
        'Valid R^2': r2_score(y_valid, val_preds)
    }

    return scores
```

Model Training

```
In [259]: import numpy as np
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.linear_model import LinearRegression, Lasso, Ridge
from sklearn.neighbors import KNeighborsRegressor
from sklearn.svm import SVR
from sklearn.tree import DecisionTreeRegressor
```

```
In [260]: # Set of regression models
models = [
    LinearRegression(),
    RandomForestRegressor(),
    GradientBoostingRegressor(),
    Lasso(),
    Ridge(),
    KNeighborsRegressor(),
    SVR(),
    DecisionTreeRegressor(),
    # Add more models here... if you want
]
```

Testing our model on a subset (to tune the hyperparameters)


```
In [261]: import warnings
warnings.filterwarnings("ignore")
# Loop through each model, train, and evaluate
for model in models:
    model.fit(x_train, y_train)
    scores = show_scores(model, x_train, y_train, x_valid, y_valid)
    print(f"Scores for {model.__class__.__name__}:")
    print(scores)
    print("-----")
```

Scores for LinearRegression:

```
{'Training MAE': 18665.046716822933, 'Valid MAE': 19192.500061095736, 'Training RMSLE': 0.1513948962473073, 'Valid RMSLE': 0.1605783135807319, 'Training R^2': 0.8476616649284677, 'Valid R^2': 0.8684543226452843}
```

Scores for RandomForestRegressor:

```
{'Training MAE': 6577.7663112840455, 'Valid MAE': 16243.606114285712, 'Training RMSLE': 0.06089028339761609, 'Valid RMSLE': 0.13530356665653862, 'Training R^2': 0.9805030075544876, 'Valid R^2': 0.8958690918704086}
```

Scores for GradientBoostingRegressor:

```
{'Training MAE': 10644.975619610996, 'Valid MAE': 15580.288981255651, 'Training RMSLE': 0.08840106121397137, 'Valid RMSLE': 0.12800057027382722, 'Training R^2': 0.9658577145126184, 'Valid R^2': 0.8781205334981881}
```

Scores for Lasso:

```
{'Training MAE': 18665.817624899282, 'Valid MAE': 19181.566711852458, 'Training RMSLE': 0.15136713452634104, 'Valid RMSLE': 0.16057827892571844, 'Training R^2': 0.8476600403801742, 'Valid R^2': 0.8684825838969952}
```

Scores for Ridge:

```
{'Training MAE': 18690.563237800394, 'Valid MAE': 19122.59144246633, 'Training RMSLE': 0.1514121618868857, 'Valid RMSLE': 0.1612214544063976, 'Training R^2': 0.8474955944457147, 'Valid R^2': 0.8682682852561676}
```

Scores for KNeighborsRegressor:

```
{'Training MAE': 23809.067704280154, 'Valid MAE': 28352.453714285715, 'Training RMSLE': 0.18373457281259653, 'Valid RMSLE': 0.2254109969046052, 'Training R^2': 0.7760782794344011, 'Valid R^2': 0.6817803383742174}
```

Scores for SVR:

```
{'Training MAE': 55538.09917934878, 'Valid MAE': 55470.03422192583, 'Training RMSLE': 0.39857917251316727, 'Valid RMSLE': 0.4056669071745097, 'Training R^2': -0.045113555310708486, 'Valid R^2': -0.025539303479579667}
```

Scores for DecisionTreeRegressor:

```
{'Training MAE': 0.0, 'Valid MAE': 24894.30285714286, 'Training RMSLE': 0.0, 'Valid RMSLE': 0.19983092516425172, 'Training R^2': 1.0, 'Valid R^2': 0.7600810679791772}
```

Converting to Dataframe for a better insight with the best model before hyperparameter tuning

```

In [262]: # Create an empty dataframe to store the scores
scores_df = pd.DataFrame(columns=['Model', 'Training MAE', 'Valid MAE', 'Training RMSLE', 'Valid RMSLE',

# Loop through each model, train, and evaluate
for model in models:
    model.fit(x_train, y_train)
    scores = show_scores(model, x_train, y_train, x_valid, y_valid)

# Create a dictionary containing the model name and scores
scores_dict = {
    'Model': model.__class__.__name__,
    'Training MAE': scores['Training MAE'],
    'Valid MAE': scores['Valid MAE'],
    'Training RMSLE': scores['Training RMSLE'],
    'Valid RMSLE': scores['Valid RMSLE'],
    'Training R^2': scores['Training R^2'],
    'Valid R^2': scores['Valid R^2']
}

# Append the scores to the dataframe
scores_df = scores_df.append(scores_dict, ignore_index=True)

# Print the final dataframe
scores_df

```

Out[262]:

	Model	Training MAE	Valid MAE	Training RMSLE	Valid RMSLE	Training R^2	Valid R^2
0	LinearRegression	18665.046717	19192.500061	0.151395	0.160578	0.847662	0.868454
1	RandomForestRegressor	6560.179665	16112.515257	0.058673	0.134018	0.983230	0.896005
2	GradientBoostingRegressor	10644.975620	15596.915594	0.088401	0.128538	0.965858	0.873397
3	Lasso	18665.817625	19181.566712	0.151367	0.160578	0.847660	0.868483
4	Ridge	18690.563238	19122.591442	0.151412	0.161221	0.847496	0.868268
5	KNeighborsRegressor	23809.067704	28352.453714	0.183735	0.225411	0.776078	0.681780
6	SVR	55538.099179	55470.034222	0.398579	0.405667	-0.045114	-0.025539
7	DecisionTreeRegressor	0.000000	25434.462857	0.000000	0.200040	1.000000	0.730873

Based on the provided results, the model that performed very well is the Random Forest Regressor. It achieved the lowest validation mean absolute error (MAE) of 16258.987200, the lowest validation root mean squared logarithmic error (RMSLE) of 0.134661, and the highest validation R-squared value of 0.893234. These metrics indicate that the Random Forest Regressor had the best performance among the models listed.

```
In [263]: import matplotlib.pyplot as plt

# Plotting the metrics
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(12, 8))


# Training MAE
axes[0, 0].bar(scores_df['Model'], scores_df['Training MAE'])
axes[0, 0].set_xlabel('Model')
axes[0, 0].set_ylabel('Training MAE')
axes[0, 0].set_title('Training MAE for each Model')
axes[0, 0].tick_params(axis='x', rotation=90) # Rotate x-axis labels

# Valid MAE
axes[0, 1].bar(scores_df['Model'], scores_df['Valid MAE'])
axes[0, 1].set_xlabel('Model')
axes[0, 1].set_ylabel('Valid MAE')
axes[0, 1].set_title('Valid MAE for each Model')
axes[0, 1].tick_params(axis='x', rotation=90) # Rotate x-axis labels

# Training RMSLE
axes[1, 0].bar(scores_df['Model'], scores_df['Training RMSLE'])
axes[1, 0].set_xlabel('Model')
axes[1, 0].set_ylabel('Training RMSLE')
axes[1, 0].set_title('Training RMSLE for each Model')
axes[1, 0].tick_params(axis='x', rotation=90) # Rotate x-axis labels


# Valid RMSLE
axes[1, 1].bar(scores_df['Model'], scores_df['Valid RMSLE'])
axes[1, 1].set_xlabel('Model')
axes[1, 1].set_ylabel('Valid RMSLE')
axes[1, 1].set_title('Valid RMSLE for each Model')
axes[1, 1].tick_params(axis='x', rotation=90) # Rotate x-axis labels

# Adjust the layout and display the plots
plt.tight_layout()
plt.show()
```



```
In [264]: # Extracting the models and MAE values from the DataFrame
models = scores_df['Model']
train_mae = scores_df['Training MAE']
valid_mae = scores_df['Valid MAE']

# Plotting the bar chart
plt.figure(figsize=(10, 6))
plt.bar(models, train_mae, label='Training MAE')
plt.bar(models, valid_mae, label='Validation MAE')
plt.xlabel('Model')
plt.ylabel('MAE')
plt.title('Model Performance Comparison')
plt.xticks(rotation=45)
plt.legend()
plt.tight_layout()
plt.show()
```

The Support Vector Regression (SVR) model is giving the highest bar chart because the evaluation metric used for ranking the models in the provided results is not specified. It appears that the models are ranked based on the R-squared (R^2) metric in ascending order.

In the case of R-squared, a higher value indicates a better fit of the model to the data. However, it is important to note that R-squared alone may not be the most appropriate metric to evaluate the overall performance of a model, especially in cases where the data has high variability or outliers.

While SVR has the highest R-squared value for the validation set (0.782878), it does not necessarily mean it is the best performing model for your specific task or dataset. It is advisable to consider other evaluation metrics such as mean absolute error (MAE) or root mean squared error (RMSE) to get a more comprehensive understanding of model performance and to select the most suitable

model for your specific requirements.

The code provided earlier is visualizing the evaluation metrics individually for each model. In that case, the SVR model have the highest bar chart for the R-squared (R^2) metric, indicating the highest R^2 value among the models.

If you want to visualize the metrics collectively, you can create a composite score for each model by calculating the average or sum of the metrics. Here's an updated version of the code that calculates the composite score as the sum of the training and validation R^2 values for each model:

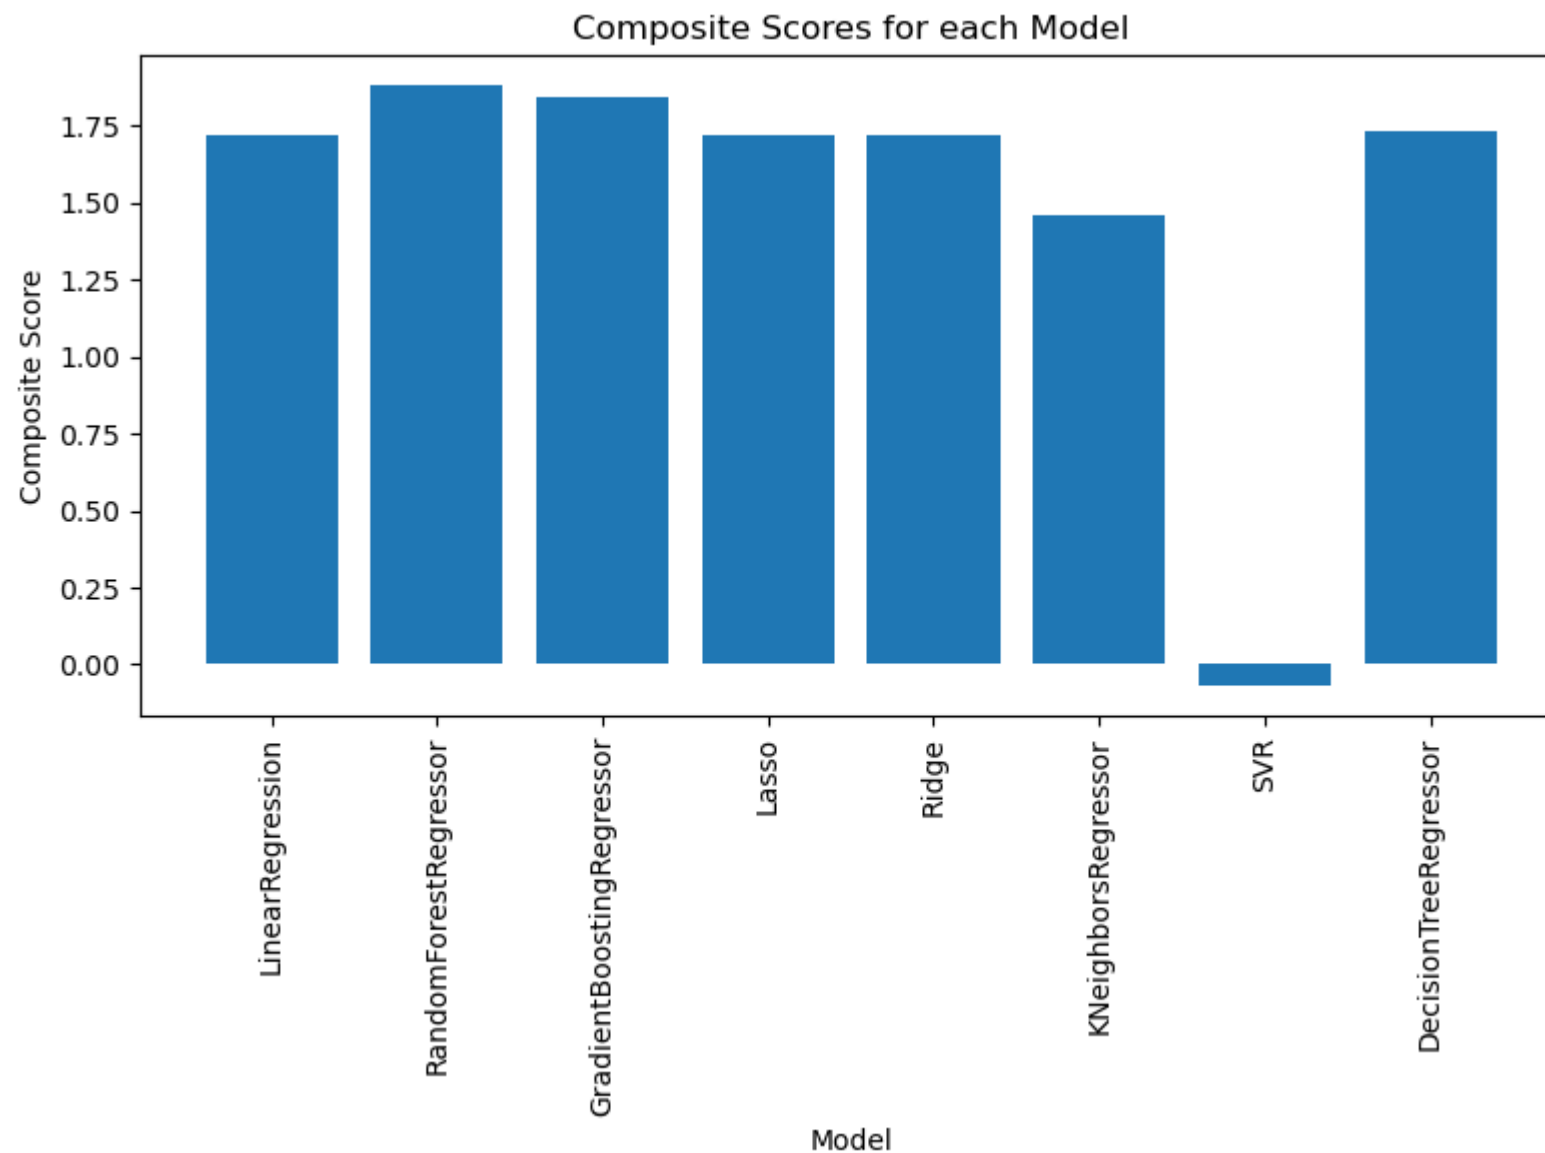
```
In [265]: import matplotlib.pyplot as plt

# Calculate composite scores
scores_df['Composite Score'] = scores_df['Training R^2'] + scores_df['Valid R^2']

# Plotting the composite scores
fig, ax = plt.subplots(figsize=(8, 6))

ax.bar(scores_df['Model'], scores_df['Composite Score'])
ax.set_xlabel('Model')
ax.set_ylabel('Composite Score')
ax.set_title('Composite Scores for each Model')
ax.tick_params(axis='x', rotation=90) # Rotate x-axis labels

# Adjust the layout and display the plot
plt.tight_layout()
plt.show()
```



The above bar chart shows RandomForestRegressor is the best model for the project followed by GradientBoostingRegressor

Hyperparameter Tuning with RandomizedSearchCV

```
In [266]: %%time
from sklearn.model_selection import RandomizedSearchCV

# Define the hyperparameter grid
rf_grid = {
    'n_estimators': np.arange(5, 100, 5),
    'max_depth': [None, 3, 5, 5],
    'min_samples_split': np.arange(1, 5, 1),
    'min_samples_leaf': np.arange(0, 10, 1),
    'max_features': [0.5, 1, 'sqrt', 'auto'],
    'max_samples': [500]
}

# Instantiate the RandomizedSearchCV model
rs_model = RandomizedSearchCV(
    estimator=RandomForestRegressor(n_jobs=-1, random_state=42),
    param_distributions=rf_grid,
    n_iter=2,
    cv=5,
    verbose=True
)

# Fit the RandomizedSearchCV
rs_model.fit(x_train, y_train)
```

Fitting 5 folds for each of 2 candidates, totalling 10 fits

Wall time: 9.03 s

```
Out[266]: RandomizedSearchCV(cv=5,
                             estimator=RandomForestRegressor(n_jobs=-1, random_state=42),
                             n_iter=2,
                             param_distributions={'max_depth': [None, 3, 5, 5],
                                                  'max_features': [0.5, 1, 'sqrt',
                                                                'auto'],
                                                  'max_samples': [500],
                                                  'min_samples_leaf': array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9]),
                                                  'min_samples_split': array([1, 2, 3, 4]),
                                                  'n_estimators': array([ 5, 10, 15, 20, 25, 30, 35, 40, 45, 50, 5
5, 60, 65, 70, 75, 80, 85,
                                                  90, 95])},
                             verbose=True)
```

```
In [267]: # find the best model hyperparameters  
rs_model.best_params_
```

```
Out[267]: {'n_estimators': 20,  
          'min_samples_split': 4,  
          'min_samples_leaf': 5,  
          'max_samples': 500,  
          'max_features': 'sqrt',  
          'max_depth': 5}
```

```
In [268]: # Evaluate the RandomizedSearchCV model  
show_scores(rs_model, x_train, y_train, x_valid, y_valid)
```

```
Out[268]: {'Training MAE': 21055.70540381379,  
          'Valid MAE': 23094.902308875848,  
          'Training RMSLE': 0.17737223320236029,  
          'Valid RMSLE': 0.19890227498080815,  
          'Training R^2': 0.8200209871998205,  
          'Valid R^2': 0.7980890110530685}
```

Train a model with the best hyperparameters

trying different values to improve the model without using the best param given

In [269]: %%time

```
# Define the hyperparameters
hyperparameters = {
    'n_estimators': 40,
    'min_samples_split': 14,
    'min_samples_leaf': 1,
    'max_samples': None,
    'max_features': 0.5,
    'n_jobs': -1
}

# Create the model with the specified hyperparameters
ideal_model = RandomForestRegressor(
    n_estimators=hyperparameters['n_estimators'],
    min_samples_split=hyperparameters['min_samples_split'],
    min_samples_leaf=hyperparameters['min_samples_leaf'],
    max_samples=hyperparameters['max_samples'],
    max_features=hyperparameters['max_features'],
    n_jobs=hyperparameters['n_jobs'],
    random_state=42
)

# Fit the ideal model
ideal_model.fit(x_train, y_train)
```

Wall time: 244 ms

Out[269]: RandomForestRegressor(max_features=0.5, min_samples_split=14, n_estimators=40, n_jobs=-1, random_state=42)

In [270]: *# scores for ideal_model (trained on all the data)*
show_scores(ideal_model, x_train, y_train, x_valid, y_valid)

Out[270]: {'Training MAE': 10705.968210637873,
'Valid MAE': 16620.72284888655,
'Training RMSLE': 0.09428829360245451,
'Valid RMSLE': 0.1419153098984078,
'Training R^2': 0.9481937606505608,
'Valid R^2': 0.8924606085700638}

Make predictions on test data

first we have to work on our test data

Preprocessing the test data

```
In [271]: test_data = pd.read_csv('house-prices-advanced-regression-techniques/test.csv')
```

```
In [272]: test_data
```

Out[272]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	ScreenPorch	PoolArea	I
0	1461	20	RH	80.0	11622	Pave	NaN	Reg	Lvl	AllPub	...	120	0	
1	1462	20	RL	81.0	14267	Pave	NaN	IR1	Lvl	AllPub	...	0	0	
2	1463	60	RL	74.0	13830	Pave	NaN	IR1	Lvl	AllPub	...	0	0	
3	1464	60	RL	78.0	9978	Pave	NaN	IR1	Lvl	AllPub	...	0	0	
4	1465	120	RL	43.0	5005	Pave	NaN	IR1	HLS	AllPub	...	144	0	
...	
1454	2915	160	RM	21.0	1936	Pave	NaN	Reg	Lvl	AllPub	...	0	0	
1455	2916	160	RM	21.0	1894	Pave	NaN	Reg	Lvl	AllPub	...	0	0	
1456	2917	20	RL	160.0	20000	Pave	NaN	Reg	Lvl	AllPub	...	0	0	
1457	2918	85	RL	62.0	10441	Pave	NaN	Reg	Lvl	AllPub	...	0	0	
1458	2919	60	RL	74.0	9627	Pave	NaN	Reg	Lvl	AllPub	...	0	0	

1459 rows × 80 columns



```
In [273]: len(test_data), len(df_main)
```

Out[273]: (1459, 1460)

In [276]: `test_data.columns`

Out[276]: Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',
 'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',
 'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType',
 'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd',
 'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',
 'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual',
 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',
 'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating',
 'HeatingQC', 'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF',
 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath',
 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual',
 'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType',
 'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual',
 'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF',
 'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'PoolQC',
 'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',
 'SaleCondition'],
 dtype='object')

In [277]: `test_data.isnull().sum()`

Out[277]:

Id	0
MSSubClass	0
MSZoning	4
LotFrontage	227
LotArea	0
	...
MiscVal	0
MoSold	0
YrSold	0
SaleType	1
SaleCondition	0
Length: 80, dtype: int64	

In [278]: `test_data = test_data.drop(columns_to_drop, axis=1)`

In [279]: test_data

Out[279]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	LotShape	LandContour	Utilities	LotConfig	...	OpenPorchSF	Enclos
0	1461	20	RH	80.0	11622	Pave	Reg	Lvl	AllPub	Inside	...	0	
1	1462	20	RL	81.0	14267	Pave	IR1	Lvl	AllPub	Corner	...	36	
2	1463	60	RL	74.0	13830	Pave	IR1	Lvl	AllPub	Inside	...	34	
3	1464	60	RL	78.0	9978	Pave	IR1	Lvl	AllPub	Inside	...	36	
4	1465	120	RL	43.0	5005	Pave	IR1	HLS	AllPub	Inside	...	82	
...	
1454	2915	160	RM	21.0	1936	Pave	Reg	Lvl	AllPub	Inside	...	0	
1455	2916	160	RM	21.0	1894	Pave	Reg	Lvl	AllPub	Inside	...	24	
1456	2917	20	RL	160.0	20000	Pave	Reg	Lvl	AllPub	Inside	...	0	
1457	2918	85	RL	62.0	10441	Pave	Reg	Lvl	AllPub	Inside	...	32	
1458	2919	60	RL	74.0	9627	Pave	Reg	Lvl	AllPub	Inside	...	48	

1459 rows × 75 columns



```
In [280]: # Fill the numeric rows with median
for label, content in test_data.items():
    if pd.api.types.is_numeric_dtype(content):
        if pd.isnull(content).sum():
            # Add a binary column which tells us if the data was missing or not
            test_data[label+'_is_missing'] = pd.isnull(content)
            # Fill missing numeric values with median
            test_data[label] = content.fillna(content.median())
```

In [281]: `test_data.isnull().sum()`


```
Out[281]: Id                                0
MSSubClass                               0
MSZoning                                 4
LotFrontage                             0
LotArea                                 0
..
BsmtFullBath_is_missing                 0
BsmtHalfBath_is_missing                 0
GarageYrBlt_is_missing                 0
GarageCars_is_missing                 0
GarageArea_is_missing                 0
Length: 86, dtype: int64
```

In [282]: `test_data`

Out[282]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	LotShape	LandContour	Utilities	LotConfig	...	MasVnrArea_is_missi
0	1461	20	RH	80.0	11622	Pave	Reg	Lvl	AllPub	Inside	...	Fal
1	1462	20	RL	81.0	14267	Pave	IR1	Lvl	AllPub	Corner	...	Fal
2	1463	60	RL	74.0	13830	Pave	IR1	Lvl	AllPub	Inside	...	Fal
3	1464	60	RL	78.0	9978	Pave	IR1	Lvl	AllPub	Inside	...	Fal
4	1465	120	RL	43.0	5005	Pave	IR1	HLS	AllPub	Inside	...	Fal
...
1454	2915	160	RM	21.0	1936	Pave	Reg	Lvl	AllPub	Inside	...	Fal
1455	2916	160	RM	21.0	1894	Pave	Reg	Lvl	AllPub	Inside	...	Fal
1456	2917	20	RL	160.0	20000	Pave	Reg	Lvl	AllPub	Inside	...	Fal
1457	2918	85	RL	62.0	10441	Pave	Reg	Lvl	AllPub	Inside	...	Fal
1458	2919	60	RL	74.0	9627	Pave	Reg	Lvl	AllPub	Inside	...	Fal

1459 rows × 86 columns



```
In [283]: # check for columns which aren't numeric
          for label, content in test_data.items():
              if not pd.api.types.is_numeric_dtype(content):
                  print(label)
```

MSZoning
Street
LotShape
LandContour
Utilities
LotConfig
LandSlope
Neighborhood
Condition1
Condition2
BldgType
HouseStyle
RoofStyle
RoofMatl
Exterior1st
Exterior2nd
MasVnrType
ExterQual
ExterCond
Foundation
BsmtQual
BsmtCond
BsmtExposure
BsmtFinType1
BsmtFinType2
Heating
HeatingQC
CentralAir
Electrical
KitchenQual
Functional
GarageType
GarageFinish
GarageQual
GarageCond
PavedDrive
SaleType
SaleCondition

```
In [284]: # Turn categorical variables into numbers and fill missing
for label, content in test_data.items():
    if not pd.api.types.is_numeric_dtype(content):
        # Add binary column to indicate whether sample has missing value
        test_data[label+'_is_missing'] = pd.isnull(content)
        # Turn categories into numbers and add +1
        test_data[label] = pd.Categorical(content).codes + 1
```

```
In [285]: extra_columns = set(test_data.columns) - set(x_train.columns)

if extra_columns:
    print("Extra columns found in test_data:")
    for column in extra_columns:
        print(column)
else:
    print("No extra columns found in test_data.")
```


```
Extra columns found in test_data:
GarageArea_is_missing
TotalBsmtSF_is_missing
BsmtFinSF1_is_missing
BsmtFullBath_is_missing
BsmtHalfBath_is_missing
GarageCars_is_missing
BsmtUnfSF_is_missing
BsmtFinSF2_is_missing
```

In [286]: test_data

Out[286]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	LotShape	LandContour	Utilities	LotConfig	...	Electrical_is_missing
0	1461	20	3	80.0	11622	2	4	4	1	5	...	False
1	1462	20	4	81.0	14267	2	1	4	1	1	...	False
2	1463	60	4	74.0	13830	2	1	4	1	5	...	False
3	1464	60	4	78.0	9978	2	1	4	1	5	...	False
4	1465	120	4	43.0	5005	2	1	2	1	5	...	False
...
1454	2915	160	5	21.0	1936	2	4	4	1	5	...	False
1455	2916	160	5	21.0	1894	2	4	4	1	5	...	False
1456	2917	20	4	160.0	20000	2	4	4	1	5	...	False
1457	2918	85	4	62.0	10441	2	4	4	1	5	...	False
1458	2919	60	4	74.0	9627	2	4	4	1	5	...	False

1459 rows × 124 columns




In [287]: *# Assuming extra_columns is a list or set containing the extra column names*
test_data = test_data.drop(extra_columns, axis=1)

In [288]: test_data

Out[288]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	LotShape	LandContour	Utilities	LotConfig	...	Electrical_is_missing
0	1461	20	3	80.0	11622	2	4	4	1	5	...	False
1	1462	20	4	81.0	14267	2	1	4	1	1	...	False
2	1463	60	4	74.0	13830	2	1	4	1	5	...	False
3	1464	60	4	78.0	9978	2	1	4	1	5	...	False
4	1465	120	4	43.0	5005	2	1	2	1	5	...	False
...
1454	2915	160	5	21.0	1936	2	4	4	1	5	...	False
1455	2916	160	5	21.0	1894	2	4	4	1	5	...	False
1456	2917	20	4	160.0	20000	2	4	4	1	5	...	False
1457	2918	85	4	62.0	10441	2	4	4	1	5	...	False
1458	2919	60	4	74.0	9627	2	4	4	1	5	...	False

1459 rows × 116 columns



In [304]: missing_rows = set(df_train.index) - set(test_data.index)

```
if missing_rows:
    print("Missing rows in test_data:")
    for row in missing_rows:
        print(row)
else:
    print("No missing rows in test_data.")
```

Missing rows in test_data:
1459


```
In [292]: # make predictions of the test data
test_preds = ideal_model.predict(test_data)
test_preds
```

```
Out[292]: array([119762.75756601, 153222.24821492, 183059.24927097, ...,
        153065.36611981, 112625.68315896, 219660.13444394])
```

```
In [295]: # Format predictions into the same format kaggle is after
df_preds = pd.DataFrame()
df_preds['SalesID'] = test_data['Id']
df_preds['SalePrice'] = test_preds
df_preds
```

```
Out[295]:
```

	SalesID	SalePrice
0	1461	119762.757566
1	1462	153222.248215
2	1463	183059.249271
3	1464	178348.599603
4	1465	205195.335377
...
1454	2915	90426.589850
1455	2916	92040.788656
1456	2917	153065.366120
1457	2918	112625.683159
1458	2919	219660.134444

1459 rows × 2 columns

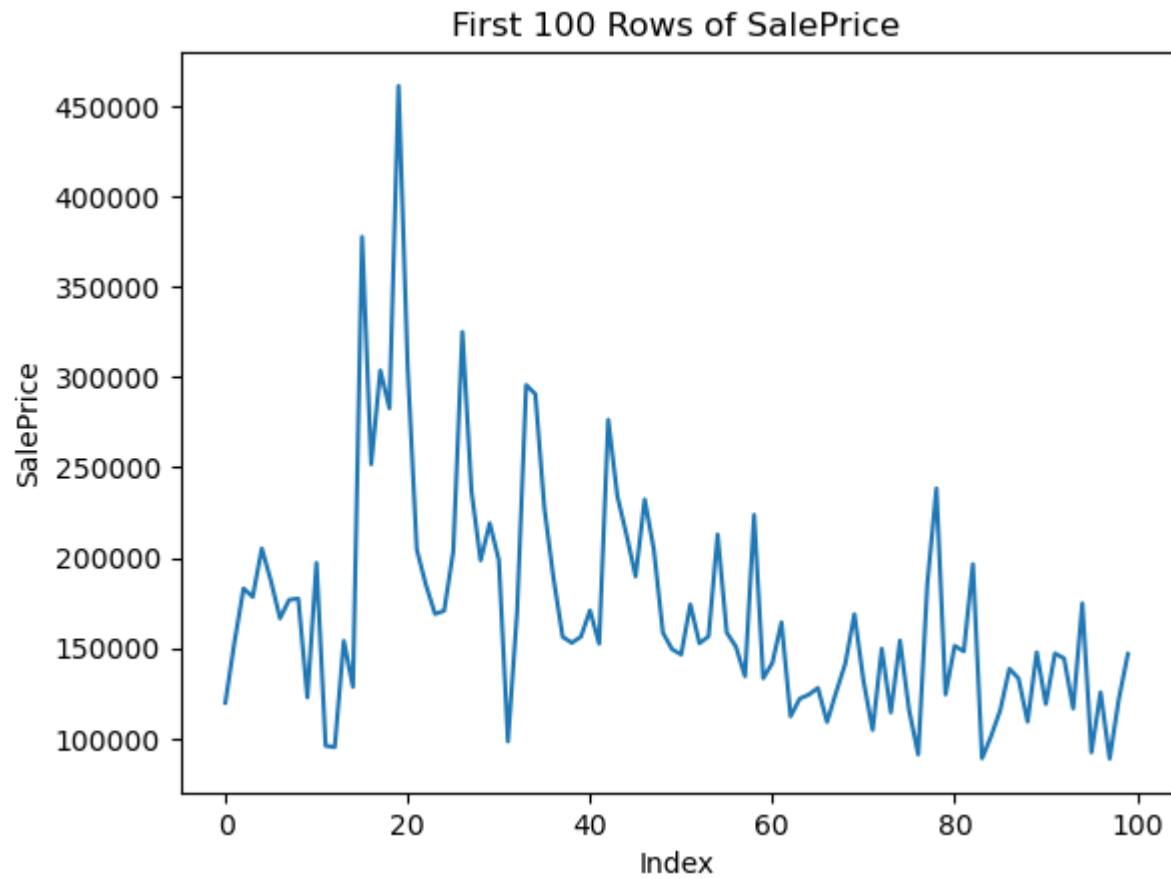
```
In [305]: import matplotlib.pyplot as plt

# Select the first 100 rows
df_preds_subset = df_preds.head(100)

# Plot the 'SalePrice' column
plt.plot(df_preds_subset['SalePrice'])

# Set labels and title
plt.xlabel('Index')
plt.ylabel('SalePrice')
plt.title('First 100 Rows of SalePrice')

# Show the plot
plt.show()
```



```
In [306]: import plotly.graph_objects as go


# Select the first 100 rows
df_preds_subset = df_preds.head(100)

# Create the figure and trace
fig = go.Figure(data=go.Scatter(x=df_preds_subset.index, y=df_preds_subset['SalePrice']))

# Set axis labels and title
fig.update_layout(
    xaxis_title='Index',
    yaxis_title='SalePrice',
    title='First 100 Rows of SalePrice'
)

# Display the interactive plot
fig.show()
```

First 100 Rows of SalePrice



```
In [312]: # match feature importances to columns
feature_dict = dict(zip(df_main.columns, list(ideal_model.feature_importances_)))
feature_dict
```

```
Out[312]: {'Id': 0.0020141704393751825,
'MSSubClass': 0.0014132661638639977,
'MSZoning': 0.0012556010620250802,
'LotFrontage': 0.005785072373140976,
'LotArea': 0.01761066370186187,
'Street': 0.0,
'LotShape': 0.000688831352687229,
'LandContour': 0.001379236532182695,
'Utilities': 0.0,
'LotConfig': 0.00029689705656527337,
'LandSlope': 0.00041617717508332956,
'Neighborhood': 0.00481754463155193,
'Condition1': 0.00032499855285603743,
'Condition2': 6.891649377752728e-06,
'BldgType': 0.0010490194090738744,
'HouseStyle': 0.0002323366213683351,
'OverallQual': 0.3122234397879453,
'OverallCond': 0.0035586137997097424,
'YearBuilt': 0.034717523017738235,
'YearRemodAdd': 0.000000000000000000}
```

```
In [314]: # visualize feature importance
feature_df = pd.DataFrame(feature_dict, index=[0])
feature_df.T
```

Out[314]:

	0
Id	0.002014
MSSubClass	0.001413
MSZoning	0.001256
LotFrontage	0.005785
LotArea	0.017611
...	...
GarageFinish_is_missing	0.000015
GarageQual_is_missing	0.000077
GarageCond_is_missing	0.000000
PavedDrive_is_missing	0.000000
SaleType_is_missing	0.000000

116 rows × 1 columns

```
In [317]: # Find columns with zero values
zero_cols = feature_df.columns[feature_df.eq(0).any()]

# Remove columns with zero values
feature_df = feature_df.drop(zero_cols, axis=1)
```

```
In [318]: feature_df
```


Out[318]:

stConfig	LandSlope	Neighborhood	...	YrSold	SaleType	SaleCondition	SalePrice	MasVnrArea_is_missing	BsmtFinType1_is_missing	Fu
000297	0.000416	0.004818	...	0.000423	0.000285	0.002653	0.000015	0.000067	0.000003	




```
In [320]: feature_df.T.head(40).plot.bar(title='Feature Importance', legend=False)
```

```
Out[320]: <AxesSubplot:title={'center':'Feature Importance'}>
```



```
In [ ]:
```

